REVEAL: Agentic In-Hand Manipulation for Language-Guided Target Presentation

Anonymous CVPR submission

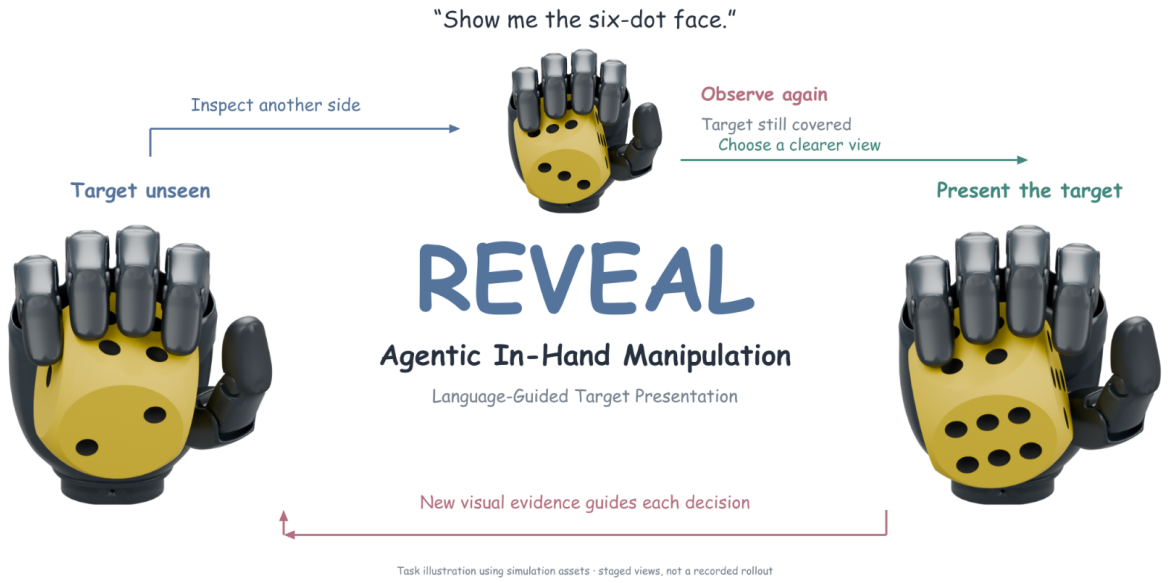


Figure 1. **REVEAL overview.** New visual evidence guides semantic decisions as the hand inspects and presents a requested object face. This staged simulation illustration is not a recorded rollout or physical-robot result.

Abstract

How can a robot find and present a requested object face when it is initially hidden or blocked by the fingers? We present REVEAL, a framework for agentic in-hand manipulation that connects visual search to continuous motor control. The agent chooses inspection actions, links the requested face to a motor goal, and directs its presentation. A recurrent motor policy follows this goal, while teacher-supervised local controllers adjust joint commands during agent inference and action execution. We evaluate six object-specific systems in simulation and test English and Chinese instructions on three objects. Under changed target-to-skill mappings, REVEAL achieves 65.8% success averaged across objects, exceeding the strongest evaluated baseline by 40.4 percentage points. Paired cube experiments show that fresh images enable additional correct target bindings, increasing success from 9/40 to 32/40. A separate controller ablation supports joint correction during inference.

1. Introduction

A command such as "present the face marked T" tells the robot what to show. The robot still has to find T, choose a motor goal that exposes it, and move the object without dropping it. The target may initially face away from the camera or be blocked by the fingers. Moving the object can reveal the target, but also changes the grasp. The robot therefore needs to decide both what to look at next and how to keep controlling the object.

Vision-language agents can interpret the instruction and inspect camera images, but each response takes time. Meanwhile, the hand and object continue to move. Repeating the previous command may lose the grasp or leave the target hidden. Slowing down limits motion but cannot change the direction of the joint command. A local controller must therefore respond to new measurements while the agent is still reasoning.

REVEAL lets the agent and motor policy run at different rates while sharing the same motor goals. The agent chooses inspection actions from the current image and action history.

Once it identifies the requested face, it links that face to a motor goal, a step we call *target binding*. A recurrent student uses this goal, measured pose and proprioception to propose joint commands. For the cube, two local controllers adjust these commands during agent inference and action execution. Both learn from teacher actions in simulation. The hand can thus respond to its current state before the next agent reply arrives.

We adapt a ViserDex-style policy [4] that estimates pose from RGB and uses recurrent control. Its new goal input accepts vectors from a fixed goal bank. The agent selects these vectors through explicit tool arguments. For comparison, a fine-tuned Qwen3-VL [2] predicts a goal vector directly, and a rule-based VLM returns a color and a readability decision. These systems share a motor student within each object, but differ in inspection, history, local control and stopping. We therefore compare complete systems and use separate cube experiments to study individual components.

Across 1,440 simulation trials on six objects, REVEAL achieves 70.4% success with the original target-to-skill mappings and 65.8% after the mappings change. These equal-object averages exceed the strongest baseline by 26.2 and 40.4 percentage points, respectively. Each object has its own policy and perception setup. A separate 2,160-trial English/Chinese study favors REVEAL over fine-tuning in all eight duck and hedgehog conditions after correction for multiple comparisons; cube differences remain inconclusive.

Paired cube studies examine how the agent acquires target evidence. Fresh images enable 25 additional correct bindings after the first call, with 23 leading to success; repeating the first image enables none. A new comparison tests adaptive against fixed inspection while keeping the model and controllers unchanged. Its 17.5-point observed success gain remains statistically unresolved. A separate three-seed study supports joint correction during agent inference.

Our contributions are threefold. First, we study language-guided target presentation through visual search, target binding and continuous in-hand control. Second, we connect an agent to a recurrent motor policy through explicit goal selection and teacher-supervised local controllers. Third, we evaluate six object-specific systems and two languages, with cube studies testing continued feedback, fresh images, inspection choices and waiting-time control.

2. Related Work

RL for dexterous in-hand manipulation. RL policies trained in simulation can reorient objects with a physical hand using domain randomization [1] and handle diverse object shapes [7]. Rapid motor adaptation handles changing object properties [15]. RotateIt combines vision, touch and proprioception through teacher–student distillation [16], while AnyRotate supports different rotation axes and hand orientations [21]. REVEAL builds on these motor capabilities to find and present a target specified by language, even when the robot must first move the object to see it.

Visual and goal-conditioned dexterity. ViserDex combines monocular pose estimation with recurrent control [4]. We adapt its released Revo3 policy to accept semantic goals. GoalDex uses current and goal images, selects goals from language with a VLM, and studies multi-object sim-to-real transfer and visual-processing delays [3]. REVEAL focuses on finding a requested face through new observations while hand control continues. ManipTrans uses residual learning for dexterous bimanual transfer [12]. Our local corrections support the agent's inspection and presentation actions.

Agentic reasoning and embodied feedback. SayCan selects language-guided skills using robotic affordances [11], and Inner Monologue uses environmental feedback to update plans [9]. Code as Policies connects perception and control APIs [14]. VoxPoser and ReKep translate language into spatial value maps and keypoint constraints [8, 10]. Magma connects multimodal representations to actions [20], while VERA includes language-guided dexterous reorientation from visual plans [13]. REVEAL studies this interaction for a target held in the hand: inspection reveals new evidence, binding selects a motor goal, and local control continues during reasoning. Separate image and inspection-policy interventions test the value of new observations and agent-selected views. Our contribution concerns their interaction with continuous in-hand control.

Asynchronous robot policies. Prior work handles inference delay through real-time action-chunk execution [5], online VLA corrections [18], latency-aware RL [22], and temporal counterfactual preferences [23]. LiMA combines sparse semantic decisions with asynchronous dexterous action generation [6]. REVEAL trains small local controllers around a motor policy and an agent with explicit tool actions. We study how observation, target binding and control work together. Asynchronous execution and residual actions alone are not new, and our ablations do not isolate the benefit of explicit controller-goal input.

3. Task and Evaluation

Semantic target. Q1 evaluates six objects in simulation: a marked cube, a wooden die, a duck, a hedgehog, a pig, and a milk carton. Q2 and Q3 use only the marked cube. Instructions request one of six surface letters, or one through six dots on the die. The cube letters are $\{T, E, R, L, S, H\}$. Each target has a fixed reference orientation $Q(g)$ that places the requested face upward. Scoring compares the full quaternion, including rotation about the face normal. Thus, showing the correct face alone is insufficient: the scorer also

requires a particular twist that the language instruction does not specify. This rule is shared by all methods.

At step t, the agent receives an RGB image I_t and instruction ℓ. The motor policy receives proprioception x_t and measured pose $\hat{p}_t$. Cube pose is estimated from RGB. Other objects use simulator pose with added noise. True pose p_t and teacher actions are used only for training and independent evaluation. During deployment, the controller cannot use a future agent reply. We denote the active goal vector by g_t, the recurrent state by h_t, and a pending request by $m_t = 1$.

Timing. Only one agent request can be pending at a time. Cube episodes allow up to six model calls and 3,600 control steps at 30 Hz. Other objects use 20 Hz over the same 120 seconds of simulated time. For the cube, a call taking L wall-clock seconds can affect control only after it completes and at least $\lceil 30L \rceil$ steps have passed since the request. Physics continues while the agent reasons. The simulator runs slower than real time, so this models inference delay but does not demonstrate a 30 Hz hardware system.

Autonomous completion. All cube systems use the same stopping check at 3 Hz. The agent must have bound the requested symbol to a goal, with no request pending and no inspection underway. Measured orientation error must be at most $\epsilon = 0.16$ radians. This check can stop the system without an explicit stop action from the agent. On other objects, REVEAL uses calibrated completion correction, while the baselines use a measured-pose check. Their stopping rules are detailed in the supplement, and all are scored with the same true-pose tolerance. At stop time τ, success is

$$S = \mathbf{1}\left[\hat{b}_\tau = b^*(\ell),\ d(q_\tau, Q(b^*)) \le \epsilon,\ \neg D_{0:\tau}\right], \quad (1)$$

where $\hat{b}_\tau$ is the declared target binding, q_τ is true orientation, $D_{0:\tau}$ indicates an earlier terminal drop, and $d(q, Q) = 2\arccos(|\langle q, Q\rangle|)$ for normalized quaternions. The target must be correct at the stop time, not merely reached earlier. We do not require an additional hold after stopping. A wrong binding or an incorrect final pose counts as failure even if the object remains in the hand. A drop terminates the trial when the object is at least 0.24 metres from its nominal grasp position.

Paired trials and service failures. Unless stated otherwise, each condition uses 40 trials: 20 per seed under two seeds. Compared methods share the same initial physical states and scored targets. Control and feedback experiments also share instructions. In the mapping experiment, the requested letter changes with the layout so that the physical target stays the same. A service error or a model call taking more than 120 seconds excludes that trial from the service-eligible comparison. We remove the same trial from every compared method, retain all raw results, and do not replace excluded trials. Each reported rate uses the denominator for its own comparison. Service failures are reported separately from control failures.

4. Method

REVEAL connects the agent to continuous motor control through a shared goal bank (Fig. 2). We first adapt a pretrained student with frozen-teacher supervision, then train local controllers around that student. At deployment, the agent selects inspection or presentation goals from new observations. We describe the cube system below; other objects use their own students, noisy pose feedback and maintenance controllers.

4.1. Visual pose estimation and motor policy

A frozen CNN takes a 120×120 RGB crop and predicts image coordinates and depth for the object center and eight keypoints. Camera calibration and rigid fitting recover object position, XYZW orientation and a fitting error. The rendered cube images include hand occlusion. The deployed policy receives this estimated pose, without ground-truth pose.

We start from a pretrained student that takes geometric goals. A learned projection replaces these goal inputs with $g_t \in \mathbb{R}^{768}$:

$$z_t = \tanh(W_g g_t + b_g) \in \mathbb{R}^{64}, \quad (2)$$

$$(u_t, h_{t+1}) = \pi_\theta([x_t, z_t], \hat{p}_t, h_t). \quad (3)$$

Here $x_t \in \mathbb{R}^{109}$ contains proprioception and history without the target pose, and $u_t \in \mathbb{R}^{21}$ contains proposed joint targets in normalized coordinates. We reuse compatible encoders, recurrent and actor weights, and normalization statistics from the released policy. The adapted student receives neither the target quaternion nor the relative-goal quaternion.

We train the goal projection and adapt the recurrent policy using actions from a frozen teacher in a DAgger-style procedure [17]. Teacher states and actions are available only during training. At deployment, the cube student uses the frozen RGB pose estimator. The supplement describes the policy and perception setup for each other object.

4.2. From observations to motor goals

The cube has a bank of six normalized goal vectors $B = \{g_1, \ldots, g_6\}$, computed from cached SigLIP2 text features [19]. Each vector represents a color skill and its reference object orientation.

Semantic agent interface. Given the instruction and current image, the hosted agent returns an explicit action: inspect, reorient, verify, stop, or abstain. We validate the action

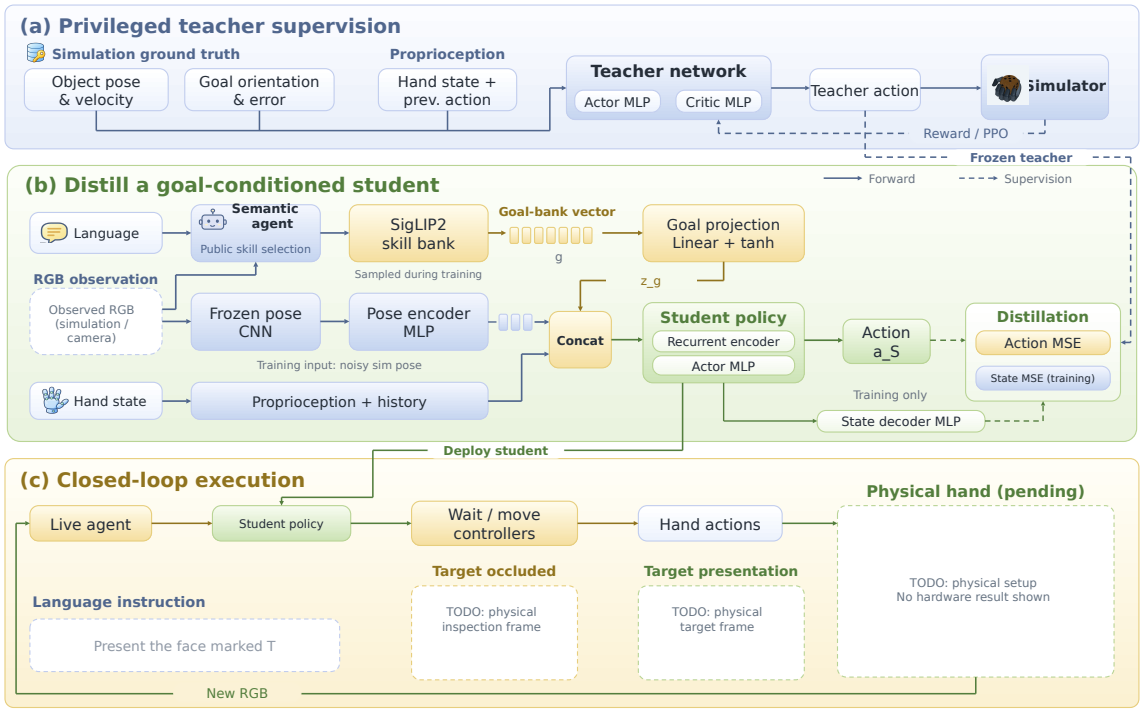


Figure 2. **REVEAL training and deployment.** Teacher supervision trains a student to follow goal-bank vectors. During deployment, the agent selects skills and local controllers adjust the hand commands. Dashed slots are reserved for future physical-hand images.

before use. Its color argument selects a goal vector from B. An inspection can move the object toward a new view without declaring that the requested target has been found. The agent chooses inspection and reorientation durations of 10–90 steps, separately from inference latency. It uses new images and the recorded action history for later decisions. Tool arguments select the student goal through a fixed lookup, without hidden agent features.

4.3. Control while the agent is reasoning

The hand needs updated commands while an agent request is pending. The waiting controller receives a feature vector s_t containing proprioception, measured pose, the student's recurrent state and proposed action, fitting error, binding status and request status. It also receives the pose and command recorded when the request began, which we call the *request anchor*, and the elapsed waiting time. The active goal g_t is encoded separately. Before target binding, this goal represents the startup or inspection skill. The controller does not know the target that a future reply may select.

A fusion layer combines the outputs of a sensor encoder and a goal encoder. The waiting actor predicts a speed logit r_t and correction values $v_t \in \mathbb{R}^{21}$. Let $\bar{q}_t$ denote normalized current joint positions and $\bar{u}_t$ the clipped student proposal.

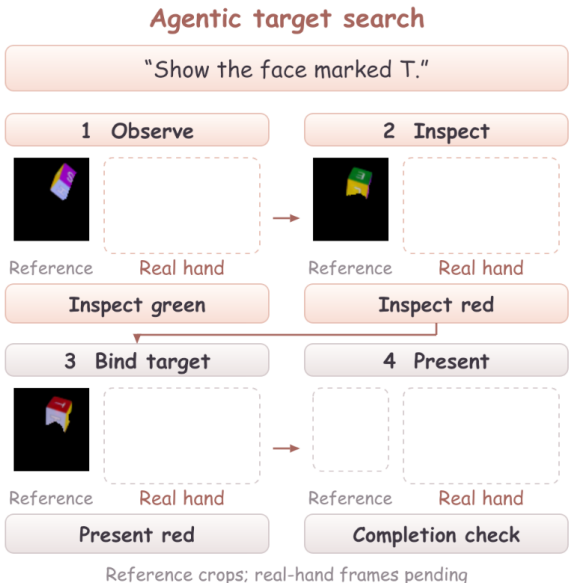


Figure 3. **Visual target search.** The agent observes, inspects candidate faces, binds the requested target, and issues a presentation command. The reference crops are simulated; real-hand frames are pending.

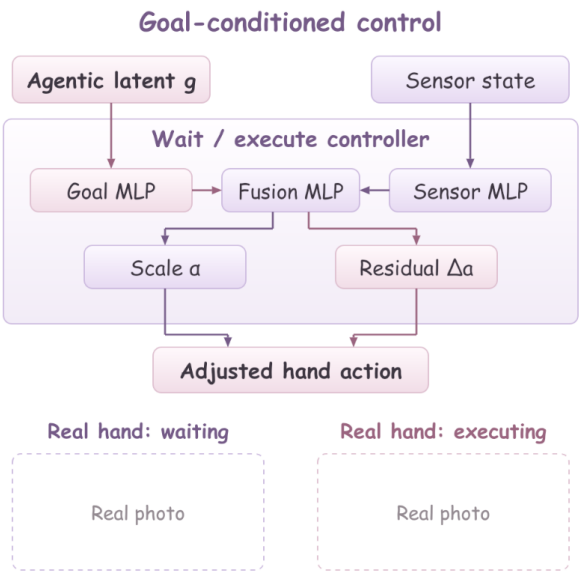


Figure 4. **Local control.** The controller combines sensor and goal features to adjust the speed and direction of the student's joint command. Empty photo slots are reserved for physical-hand results.

The applied command is

$$\alpha_t = \mathrm{sigmoid}(4r_t + \mathrm{logit}(0.15)), \tag{4}$$

$$b_t = \begin{cases} \bar{q}_t + \alpha_t(\bar{u}_t - \bar{q}_t), & m_t = 1, \\ \bar{u}_t, & m_t = 0, \end{cases} \tag{5}$$

$$a_t = \mathrm{clip}\big(b_t + 0.35\tanh(v_t), -1, 1\big). \tag{6}$$

Zero actor outputs give fixed slowing with coefficient 0.15 while a request is pending. The environment keeps the same joint mapping and smoothing. Unlike scalar slowing, the joint correction can change the command direction as well as its size.

During training, the frozen teacher provides holding actions for the request anchor. At deployment, the actor uses only measured inputs and the active goal, without the pending reply. Writing the correction as $\delta_\phi = 0.35\tanh(v_t)$, we train it to match the teacher action while keeping this correction small:

$$\mathcal{L}_{\text{control}} = \mathbb{E}_{(s,g,a^*)\sim\mathcal{D}}\left[\|a_\phi(s,g) - a^*\|_2^2 + 0.01\|\delta_\phi(s,g)\|_2^2\right]. \tag{7}$$

The waiting-actor update changes only control during pending requests. The student, pose estimator, execution controller and stopping rule remain fixed. Training details are in the supplement.

4.4. Learning execution after a decision

When no request is pending ($m_t = 0$), a separate execution actor supplies v_t in Eq. 6. It learns from teacher actions to add a bounded correction to the student's proposal. The actor also receives the active goal, although our experiments do not isolate the benefit of this extra goal input.

4.5. Closed-loop execution

At each control step, REVEAL updates the pose estimate, runs the recurrent student, selects the waiting or execution actor, and sends the resulting joint targets to the hand. Only a valid agent reply can change the active skill or target binding. The next image shows the object after the motion that occurred during inference. A separate stopping check determines whether to end the trial. Invalid replies cannot select the true target through an oracle fallback.

We evaluate student adaptation separately from controller learning because a new student also changes the recurrent features and action proposals seen by the controllers.

5. Experiments

Q1 compares complete systems, Q2 studies control during inference, and Q3 studies active target search. Q2 and Q3 use only the marked cube. Q4 describes the pending hardware evaluation.

Q1. How do complete systems compare across objects, mappings and instruction variants?
Q2. How should the hand move while the agent is reasoning?
Q3. How do new observations and inspection choices help find the target?
Q4. Does the system work on physical objects?

5.1. Setup and reporting

We use Revo3 assets in Isaac Lab 3.0/Newton. Unless specified otherwise, each condition has 40 trials across two evaluation seeds, paired by physical task. Cube systems share the frozen RGB pose estimator, goal bank, student, stopping rule and 0.16-radian tolerance. Other objects use separate students and noisy simulator pose. REVEAL additionally uses calibrated completion correction on these objects. The VLM baselines omit the extra control actors and follow a fixed inspection schedule when the target is unreadable. REVEAL chooses its inspection actions. These differences and measured inference delays are part of the complete-system comparison.

The current cube system updates only the waiting actor. Q1 and Q3 compare methods on the same service-eligible trials. Q2 reports all 120 trials per method. Q1 and Q2 use different trial sets and answer different questions, so their absolute success rates need not match. Controller gains are measured within Q2. Its three-seed table reports descriptive results. The supplement gives the paired tests and corrections for multiple comparisons used in each study.

5.2. Q1: Systems across objects and mappings

We compare the original and changed target-to-skill mappings on six objects. Earlier native-image studies are reported separately in the supplement. Neither found a significant REVEAL–fine-tuned difference (cube $p = 0.267$, die $p = 0.4545$).

Aggregate system performance. Across 1,440 scheduled trials, mean success is 70.4% with original mappings and 65.8% with changed mappings. The strongest baseline averages are 44.2% for fine-tuning and 25.4% for rules, respectively. REVEAL improves on them by 26.2 and 40.4 percentage points. These post-hoc averages weight each object's service-eligible rate equally. They do not pool trials or provide an additional significance test.

Four additional objects. Duck, hedgehog, pig and milk carton contribute 960 trials across two mappings. REVEAL has the highest observed rate in each condition (Table 1), with a separate policy for each object. Three service failures exclude matching trials from all methods. The milk carton still drops in 15/40 original-mapping and 20/40 changed-mapping REVEAL trials. Thus, higher completion does not guarantee a stable grasp. Reset checks and sensitivity analyses are in the supplement.

Visual-semantic intervention. We change only the image sent to the agent, keeping pose-estimator input and lighting fixed. Each method uses 40 matched states under three conditions: original letters, cyclically reassigned letters, and hidden letters. All three use reconstructed face images. After reassignment, REVEAL remains at 32/40 successes, fine-tuning falls from 30/40 to 16/40 (Holm $p = 0.0130$, six comparisons), and rules remain at 23/40. Hiding letters yields 0/39, 3/40 and 0/40, respectively. Nine of eleven fine-tuned wrong bindings after reassignment follow the old mapping. This suggests learned shortcuts, although it does not prove memorization as the sole cause. One reassignment cannot establish broad generalization or equivalence.

Dice layout transfer. A separate 240-trial study changes the die's visible layout by rotating its visual mesh. Collision geometry, inertia and the prompt separating dot counts from motor-skill IDs stay fixed. Each condition has 40 matched trials with no exclusions. Under the changed layout, exploratory paired tests give $p = 2.98 \times 10^{-8}$ against fine-tuning and $p = 5.96 \times 10^{-8}$ against rules. REVEAL's own layout difference is not significant ($p = 0.688$), which does not prove equivalence. REVEAL also uses inspection, history and noisy feedback, so the result compares complete systems.

Language and wording coverage. We test standard and rewritten instructions in English and Chinese on three objects: 60 trials per condition across three seeds, totaling 2,160 trials. Table 2 compares the same service-eligible trials for all methods. We keep weights, mappings and scoring fixed, without retraining for Chinese. The rule baseline also translates its task/output instructions. This tests language coverage of the deployed systems. The supplement retains all raw outcomes, including service failures.

REVEAL reaches 67.8–76.7% on duck and 73.3–76.7% on hedgehog. It outperforms fine-tuning in all eight conditions (Holm-adjusted paired $p \leq 0.0258$ across 36 method comparisons). Six of eight differences against rules remain significant after correction. On cube, REVEAL leads numerically in three conditions, while fine-tuning leads for rewritten Chinese instructions (81.1% versus 67.9%). None of these four differences is significant (all adjusted $p = 1.0$). Requiring exact state and image matching preserves these significance decisions. No language or wording comparison is significant after its separate 36-test correction, which does not prove equivalence.

5.3. Q2: Controller ablations on the cube

We compare four ways to control the hand while the agent is reasoning: continue the student's command, slow it by a fixed coefficient of 0.15, learn the slowing coefficient, or learn slowing together with a bounded joint correction. All four use the same execution controller when no request is pending. Each method has 120 trials, with 40 per seed across three seeds.

Learned slowing with joint correction achieves 104/120 successes (86.7%), compared with 88/120 (73.3%) for fixed slowing, a gain of 13.3 percentage points. Waiting drops decrease from 13 to five, and execution drops from eight to four. Learning speed alone gives 82/120 successes (68.3%), below fixed slowing. The results favor learning speed and joint correction together. Drop stages tell us when failure occurs, but do not identify its cause. Seed settings and outcome breakdowns are in the supplement.

5.4. Q3: Active Target Search and Visual Feedback

Three separate paired experiments examine the search process (Table 4): continuing to observe and decide, receiving new images, and choosing where and how long to inspect. Each comparison uses its own trial set.

(a) Does continued feedback help find the target? The initial plan selects up to six bounded actions from one image and measured feedback. Full feedback allows up to six calls with new images and history. Both use the same local controllers, stopping rule and time limit. Success rises from 6/39 to 33/39 eligible trials: 27 pairs succeed only with feedback and none only with the initial plan (exact $p =$

Object	Motor pose	$n_{\text{orig}}/n_{\text{changed}}$	Original mapping			Changed mapping		
			Rule	Fine-tuned	REVEAL	Rule	Fine-tuned	REVEAL
Cube	RGB pose	40 / 40	57.5	75.0	**80.0**	57.5	40.0	**80.0**
Die	Noisy pose	40 / 40	32.5	57.5	**70.0**	2.5	0.0	**65.0**
Duck	Noisy pose	38 / 40	39.5	36.8	**65.8**	22.5	5.0	**60.0**
Hedgehog	Noisy pose	39 / 40	43.6	33.3	**71.8**	25.0	5.0	**80.0**
Pig	Noisy pose	40 / 40	52.5	40.0	**92.5**	37.5	7.5	**82.5**
Milk carton	Noisy pose	40 / 40	32.5	22.5	**42.5**	7.5	2.5	**27.5**
Mean	Object-specific	–	43.0	44.2	**70.4**	25.4	10.0	**65.8**

Table 1. **Q1: six objects and two semantic mappings.** Completion (%) uses common service-eligible trials; each cell has 40 scheduled trials. Mean weights six objects equally. Bold denotes the highest observed rate, not significance. Object-specific weights remain fixed across mappings; bare baselines omit added control modules. Cube uses reconstructed upper images and RGB pose; die rotates its visual mesh; other objects permute surface letters with noisy simulator pose. Native-image and language cohorts remain separate.

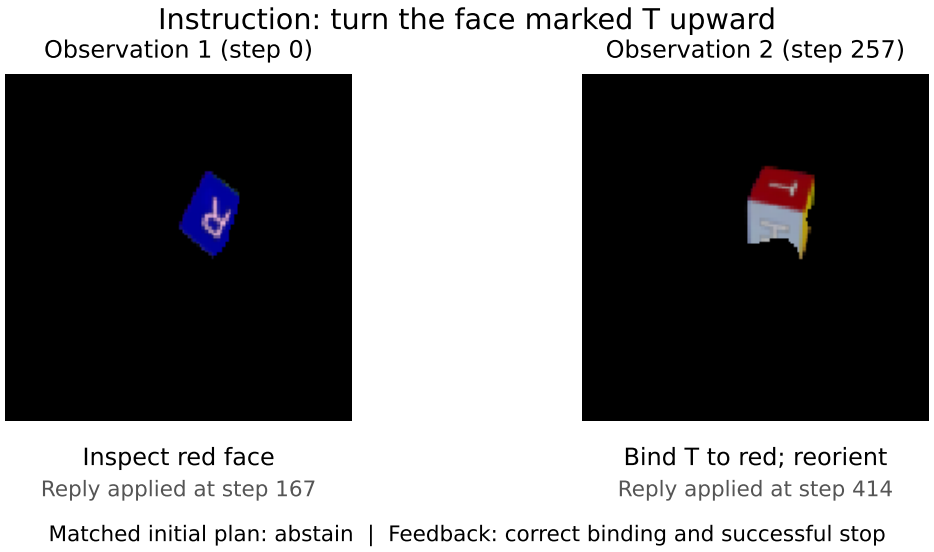


Figure 5. **Recorded inspection before binding.** Unretouched model-input frames from a recorded simulated trial. The initial plan abstains; feedback reveals T, then REVEAL binds T to red and completes at step 420 (true error 0.0812 radians). Both arms start from identical state and image. This selected example illustrates evidence acquisition, not wrong-binding correction or hardware performance.

1.49×10^{-8}). One service failure excludes a matched pair. Of 28 initial-plan abstentions, 22 succeed with feedback; 32 feedback trials first bind after inspection. Figure 5 illustrates this search process. This comparison also changes the call budget.

(b) Does the next image provide useful evidence? Both conditions allow six calls with the same request rule, live numeric feedback, history and local control. Only the agent's RGB input changes: the current image versus the exact first frame on every call. Success is 32/40 versus 9/40, with all 23 pairs with different outcomes favoring fresh images (exact $p = 2.38 \times 10^{-7}$). Correct terminal bindings number 34 versus nine. Both conditions first bind correctly on call one in the same nine trials. Fresh images enable 25 additional correct bindings later, of which 23 lead to success; repeated first frames enable none. Drops differ by only one. These records connect new observations to target binding and completion. They measure declared bindings, not initial visibility or first visual recognition.

(c) Does choosing the next view help? Forty new pairs compare agent-selected inspection directions and durations with a fixed color sweep of 90-step actions. Both use the same model, controllers, fresh images, history format and stopping rule. Feedback reports the action actually executed. The model still decides whether to inspect, bind or abstain in both conditions. Adaptive inspection achieves 38/40 successes (95.0%), compared with 31/40 (77.5%) for the fixed sweep. Nine pairs favor adaptive and two favor fixed inspec-

System	English Standard	English Rewrite	Chinese Standard	Chinese Rewrite
Cube (n)	56	59	57	53
Rule VLM	48.2	45.8	31.6	32.1
Fine-tuned VLM	71.4	72.9	75.4	**81.1**
REVEAL	**75.0**	**79.7**	**84.2**	67.9
Duck (n)	60	59	60	60
Rule VLM	50.0	25.4	33.3	31.7
Fine-tuned VLM	38.3	30.5	31.7	35.0
REVEAL	**68.3**	**67.8**	**68.3**	**76.7**
Hedgehog (n)	60	60	52	60
Rule VLM	46.7	38.3	36.5	36.7
Fine-tuned VLM	38.3	31.7	34.6	25.0
REVEAL	**73.3**	**76.7**	**75.0**	**76.7**

Table 2. **Q1: language and wording coverage.** Completion (%) on common service-eligible trials. Methods share the listed n of 60 scheduled slots. Bold denotes the highest observed rate, not significance. Raw counts and paired analyses are in the supplement.

Waiting policy	Success	%	Wait drops	Exec. drops	Other
Continue student	73/120	60.8	24	14	9
Fixed slowing	88/120	73.3	13	8	11
Learned slowing	82/120	68.3	7	14	17
Slowing + correction	**104/120**	**86.7**	**5**	**4**	**7**

Table 3. **Q2: controller ablation on the cube.** Each method has 120 trials across three seeds. Wait and execution drops denote failure stages; Other contains the remaining failures. Outcome counts sum to 120 in each row.

Condition	Success	%	Drops	Calls
(a) Continued feedback: one vs. up to six calls				
One initial plan	6/39	15.4	0	40
Full feedback	33/39	84.6	5	120
(b) Fresh visual evidence: same six-call cap				
Current image	32/40	80.0	3	112
Repeated first frame	9/40	22.5	4	98
(c) Inspection policy: same model and controllers				
Fixed sweep	31/40	77.5	1	126
Adaptive inspection	38/40	95.0	0	117

Table 4. **Q3: active target search on the cube.** Each panel has 40 scheduled trials per condition. (a) excludes one service-affected pair and reuses Q2's initial full-feedback subset. (b) changes only the agent's RGB input. (c) changes inspection direction and duration. (b,c) have no exclusions. Drops/calls include all trials; actual inference cost differs.

tion, giving a 17.5-point observed gain (exact $p = 0.0654$). Correct terminal bindings are 39 versus 32, with zero versus one drop. This suggests a benefit from choosing inspection actions, but the success difference is not significant at $\alpha = 0.05$ and tests only one fixed sweep.

5.5. Q4: Multi-object hardware evaluation

Hardware protocol. **[TODO: Complete the multi-object physical-hand evaluation.]** No hardware outcomes are available yet. The planned comparison uses REVEAL and two baselines per object, reporting completion, drops, false stops and inference cost. The supplement details calibration, resets, independent scoring and synchronized videos.

6. Discussion and Limitations

Search and control. Fresh images help the agent bind targets that it does not bind from the first view. The inspection-policy comparison shows a favorable but statistically unresolved trend. Together with the controller ablation, these studies examine evidence gathering and control during inference. Q1 compares complete systems and cannot separate all model, inspection, controller and stopping effects. Q2/Q3 mechanisms are tested only on the cube.

Evaluation limits. Two Q2 seeds are available only as combined counts, preventing seed-level variability estimates and paired tests. We have not isolated the benefit of controller-goal input or compared with a generic time-aware residual controller. Equal call limits do not ensure equal inference cost. We test one changed mapping per object and limited English/Chinese templates. Success checks binding and pose at termination, without an independent visibility check or a post-stop hold. It also requires a twist not specified by the instruction.

Deployment scope. Each object uses separate weights. Cube pose comes from rendered RGB; other objects use noisy simulator pose. Real-camera performance, transfer to new objects and hardware execution remain untested. Hosted models may change over time, limiting reproducibility. Rates excluding service failures do not measure end-to-end availability.

7. Conclusion

REVEAL connects active visual search and target binding to continuous in-hand control. Across six object-specific systems, mean success is 70.4% with original mappings and 65.8% with changed mappings. Cube studies show how fresh observations support target binding and how joint correction helps during inference. Adaptive inspection offers a promising but statistically unresolved gain over a fixed sweep. Physical-hand tests remain necessary to assess deployment.

LLM Usage Statement

OpenAI Codex assisted with research discussion, implementation, experiment orchestration, literature retrieval, statistical analysis, figure-layout code, and drafting and editing throughout the manuscript. A hosted vision-language model and Qwen3-VL are evaluated system components. Numerical claims are grounded in recorded experiments; unfinished work is identified explicitly. The authors are responsible for verifying the text, references, code, figures, and results before submission.

References

[1] Marcin Andrychowicz, Bowen Baker, Maciek Chociej, Rafal Józefowicz, Bob McGrew, Jakub Pachocki, Arthur Petron, Matthias Plappert, Glenn Powell, Alex Ray, Jonas Schneider, Szymon Sidor, Josh Tobin, Peter Welinder, Lilian Weng, and Wojciech Zaremba. Learning dexterous in-hand manipulation. *The International Journal of Robotics Research*, 39(1):3–20, 2020. 2

[2] Shuai Bai, Yuxuan Cai, Ruizhe Chen, Keqin Chen, Xionghui Chen, Zesen Cheng, Lianghao Deng, Wei Ding, Chang Gao, Chunjiang Ge, Wenbin Ge, Zhifang Guo, Qidong Huang, Jie Huang, Fei Huang, Binyuan Hui, Shutong Jiang, Zhaohai Li, Mingsheng Li, Mei Li, Kaixin Li, Zicheng Lin, Junyang Lin, Xuejing Liu, Jiawei Liu, Chenglong Liu, Yang Liu, Dayiheng Liu, Shixuan Liu, Dunjie Lu, Ruilin Luo, Chenxu Lv, Rui Men, Lingchen Meng, Xuancheng Ren, Xingzhang Ren, Sibo Song, Yuchong Sun, Jun Tang, Jianhong Tu, Jianqiang Wan, Peng Wang, Pengfei Wang, Qiuyue Wang, Yuxuan Wang, Tianbao Xie, Yiheng Xu, Haiyang Xu, Jin Xu, Zhibo Yang, Mingkun Yang, Jianxin Yang, An Yang, Bowen Yu, Fei Zhang, Hang Zhang, Xi Zhang, Bo Zheng, Humen Zhong, Jingren Zhou, Fan Zhou, Jing Zhou, Yuanzhi Zhu, and Ke Zhu. Qwen3-VL technical report. *arXiv preprint arXiv:2511.21631*, 2025. 2

[3] Arjun Bhardwaj, Mayank Mittal, Maximum Wilder-Smith, Vaishakh Patil, and Marco Hutter. GoalDex: Learning image-goal-conditioned policies for dexterous in-hand manipulation. *arXiv preprint arXiv:2611.02171*, 2026. 2

[4] Arjun Bhardwaj, Maximum Wilder-Smith, Mayank Mittal, Vaishakh Patil, and Marco Hutter. ViserDex: Visual sim-to-real for robust dexterous in-hand reorientation. *arXiv preprint arXiv:2604.11138*, 2026. 2

[5] Kevin Black, Manuel Galliker, and Sergey Levine. Real-time execution of action chunking flow policies. In *Advances in Neural Information Processing Systems*, pages 33383–33407, 2025. 2

[6] Ning Chen, Yankai Fu, Junkai Zhao, Qianpu Sun, Guocai Yao, Pengwei Wang, Zhongyuan Wang, and Shanghang Zhang. LiMA: Bridging long-term imagination to real-time dexterous manipulation via asynchronous diffusion. *arXiv preprint arXiv:2609.28431*, 2026. 2

[7] Tao Chen, Jie Xu, and Pulkit Agrawal. A system for general in-hand object re-orientation. In *Proceedings of the 5th Conference on Robot Learning*, pages 297–307, 2022. 2

[8] Wenlong Huang, Chen Wang, Ruohan Zhang, Yunzhu Li, Jiajun Wu, and Li Fei-Fei. VoxPoser: Composable 3D value maps for robotic manipulation with language models. In *Proceedings of the 7th Conference on Robot Learning*, pages 540–562, 2023. 2

[9] Wenlong Huang, Fei Xia, Ted Xiao, Harris Chan, Jacky Liang, Pete Florence, Andy Zeng, Jonathan Tompson, Igor Mordatch, Yevgen Chebotar, Pierre Sermanet, Tomas Jackson, Noah Brown, Linda Luu, Sergey Levine, Karol Hausman, and Brian Ichter. Inner monologue: Embodied reasoning through planning with language models. In *Proceedings of the 6th Conference on Robot Learning*, pages 1769–1782, 2023. 2

[10] Wenlong Huang, Chen Wang, Yunzhu Li, Ruohan Zhang, and Li Fei-Fei. ReKep: Spatio-temporal reasoning of relational keypoint constraints for robotic manipulation. In *Proceedings of the 8th Conference on Robot Learning*, pages 4573–4602, 2025. 2

[11] Brian Ichter, Anthony Brohan, Yevgen Chebotar, Chelsea Finn, Karol Hausman, Alexander Herzog, Daniel Ho, Julian Ibarz, Alex Irpan, Eric Jang, Ryan Julian, Dmitry Kalashnikov, Sergey Levine, Yao Lu, Carolina Parada, Kanishka Rao, Pierre Sermanet, Alexander T. Toshev, Vincent Vanhoucke, Fei Xia, Ted Xiao, Peng Xu, Mengyuan Yan, Noah Brown, Michael Ahn, Omar Cortes, Nicolas Sievers, Clayton Tan, Sichun Xu, Diego Reyes, Jarek Rettinghouse, Jornell Quiambao, Peter Pastor, Linda Luu, Kuang-Huei Lee, Yuheng Kuang, Sally Jesmonth, Nikhil J. Joshi, Kyle Jeffrey, Rosario Jauregui Ruano, Jasmine Hsu, Keerthana Gopalakrishnan, Byron David, Andy Zeng, and Chuyuan Kelly Fu. Do as i can, not as i say: Grounding language in robotic affordances. In *Proceedings of the 6th Conference on Robot Learning*, pages 287–318, 2023. 2

[12] Kailin Li, Puhao Li, Tengyu Liu, Yuyang Li, and Siyuan Huang. ManipTrans: Efficient dexterous bimanual manipulation transfer via residual learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 6991–7003, 2025. 2

[13] Sizhe Lester Li, Evan Kim, Xingjian Bai, Tong Zhao, Tao Pang, Max Simchowitz, and Vincent Sitzmann. Turning video models into generalist robot policies. *arXiv preprint arXiv:2605.27817*, 2026. 2

[14] Jacky Liang, Wenlong Huang, Fei Xia, Peng Xu, Karol Hausman, Brian Ichter, Pete Florence, and Andy Zeng. Code as policies: Language model programs for embodied control. In *Proceedings of the IEEE International Conference on Robotics and Automation*, pages 9493–9500, 2023. 2

[15] Haozhi Qi, Ashish Kumar, Roberto Calandra, Yi Ma, and Jitendra Malik. In-hand object rotation via rapid motor adaptation. In *Proceedings of the 6th Conference on Robot Learning*, pages 1722–1732, 2023. 2

[16] Haozhi Qi, Brent Yi, Sudharshan Suresh, Mike Lambeta, Yi Ma, Roberto Calandra, and Jitendra Malik. General in-hand object rotation with vision and touch. In *Proceedings of the 7th Conference on Robot Learning*, pages 2549–2564, 2023. 2

[17] Stephane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the Fourteenth*

International Conference on Artificial Intelligence and Statistics, pages 627–635, 2011. 3

[18] Kohei Sendai, Maxime Alvarez, Tatsuya Matsushima, Yutaka Matsuo, and Yusuke Iwasawa. Leave no observation behind: Real-time correction for VLA action chunks. *arXiv preprint arXiv:2509.23224*, 2025. 2

[19] Michael Tschannen, Alexey Gritsenko, Xiao Wang, Muhammad Ferjad Naeem, Ibrahim Alabdulmohsin, Nikhil Parthasarathy, Talfan Evans, Lucas Beyer, Ye Xia, Basil Mustafa, Olivier Hénaff, Jeremiah Harmsen, Andreas Steiner, and Xiaohua Zhai. SigLIP 2: Multilingual vision-language encoders with improved semantic understanding, localization, and dense features. *arXiv preprint arXiv:2502.14786*, 2025. 3

[20] Jianwei Yang, Reuben Tan, Qianhui Wu, Ruijie Zheng, Baolin Peng, Yongyuan Liang, Yu Gu, Mu Cai, Seonghyeon Ye, Joel Jang, Yuquan Deng, and Jianfeng Gao. Magma: A foundation model for multimodal AI agents. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14203–14214, 2025. 2

[21] Max Yang, Chenghua Lu, Alex Church, Yijiong Lin, Christopher J. Ford, Haoran Li, Efi Psomopoulou, David A. W. Barton, and Nathan F. Lepora. AnyRotate: Gravity-invariant in-hand object rotation with sim-to-real touch. In *Proceedings of the 8th Conference on Robot Learning*, pages 4727–4747, 2025. 2

[22] Brian Zhu, Momen Khalil, E Harrison, Emanuele Poggi, Philipp Schmitt, Bernd Kast, Philine Meister, Pranav Atreya, Qiyang Li, Finn Ferchau, Cesar Colmenero, Yash Shahapurkar, Gokul Narayanan, Melih Erdogan, Kai Wurm, Georg von Wichert, Oier Mees, Eugen Solowjow, Andrew Wagenmaker, and Sergey Levine. Learning to act while waiting: RL finetuning of generalist robot policies under inference latency. *arXiv preprint arXiv:2608.23831*, 2026. 2

[23] Yixiang Zhu, Yonghao Chen, Jiaxiang Zou, Yusong Hu, and Xinyu Chen. DEFLECT: Temporal counterfactual preference learning for delay-robust asynchronous VLAs. *arXiv preprint arXiv:2605.19294*, 2026. 2

REVEAL: Supplementary Material

A. Overview

This supplement accompanies the manuscript updated on 7 November 2026. Q1 compares six object-specific systems in simulation. Q2 studies control during inference; Q3 studies active target search and visual feedback on the marked cube. Q2 uses only seeds 187102, 187103 and 187104, with 40 trials per seed and method. Other seeds below belong to Q1, Q3 or training diagnostics. We also report the 120-trial native-layout die comparison, the 240-trial die layout study, 960 trials on four additional Q1 objects, and 2,160 language trials on cube, duck and hedgehog. The independent cube pose study provides supporting simulation evidence. Physical-hand tests and downstream tests of a replacement pose estimator remain unfinished.

The current complete-system comparison updates only REVEAL's waiting actor. The student, pose estimator, execution actor and success rule stay fixed. Earlier studies use their recorded controller versions and are reported separately. We exclude invalid contact-capacity runs and do not combine results from the earlier RGB-D system or different protocols. The original baseline-only comparison was added after the first study, so we do not describe the full historical evaluation as preregistered.

B. Task and Runtime

This section describes the cube system. The die, additional-object and language sections explain the differences for other objects.

B.1. Object and goals

Table S1 records the fixed face mapping. Body normals include the original mesh rotation. All methods use the same mapping and cached six-vector bank. The cyan channel corresponds visually to a pale-blue face and magenta to a purple face. The target quaternion gives a canonical twist as well as the upward face normal.

Letter	Skill color	Body normal	Target XYZW
T	red	$(0, -1, 0)$	$(-s, 0, 0, s)$
E	green	$(-1, 0, 0)$	$(0, s, 0, s)$
R	blue	$(0, 0, -1)$	$(1, 0, 0, 0)$
L	yellow	$(0, 0, 1)$	$(0, 0, 0, 1)$
S	magenta	$(0, 1, 0)$	$(s, 0, 0, s)$
H	cyan	$(1, 0, 0)$	$(0, -s, 0, s)$

Table S1. Frozen goal mapping, with $s = \sqrt{1/2}$.

The instruction names a letter, and the benchmark checks its fixed reference orientation with a 0.16-radian tolerance. A drop ends the trial when the object moves at least 0.24 metres from its nominal grasp position. This rule does not detect every brief slip or loss of contact within that boundary.

B.2. Physics and perception

The runtime uses Isaac Lab 3.0, Isaac Sim 6.2, Newton 1.6.1, Warp 1.17, and PyTorch 2.11 with CUDA 13.0. The validated collision model contains 23 bodies and 117 shapes, comprising 116 convex shapes and one box. The contact allocation uses `nconmax=4096` and `njmax=1024`. Earlier overflow runs are invalid even if their reported rewards appear plausible.

The pose estimator is frozen. It predicts 27 image-coordinate/depth outputs for nine points, then recovers pose through calibrated geometry. The object image is rendered at 120×120 with live hand occlusion. Sensor calibration, object asset, policy normalization, joint order, action mapping, and smoothing belong to the release contract. A change to one of these produces a new release. The current setup does not establish real-camera robustness or unseen-scene generalization.

B.3. Timing, termination, and exclusions

Each episode lasts at most 120 seconds of simulated control time. A model call taking L wall-clock seconds must remain pending for at least $\lceil 30L \rceil$ control steps, and its reply cannot be used before the call finishes. Physics continues throughout. Since simulation runs slower than real time, an episode takes longer in wall-clock time. Hardware control speed must be measured separately.

The stopping check runs at 3 Hz. The requested symbol must be bound to a goal, with no inspection or model request underway, and estimated orientation error must be at most 0.16 radians. The independent scorer then checks the actual requested face and true orientation using the same tolerance. It scores the stop time, without an extra holding period or credit for a better pose reached earlier.

A service error or model call exceeding 120 seconds is recorded without replacement. The original three-system panel removes trial 183102011 from every arm. This leaves 39 common eligible trials for that historical Q1 panel. The current Q2 study reports all 120 outcomes per method. Raw 40-trial outcomes remain visible, including service-affected cases.

C. Models and Training

Figure S1 expands the main-text task overview with the training and deployment pathways.

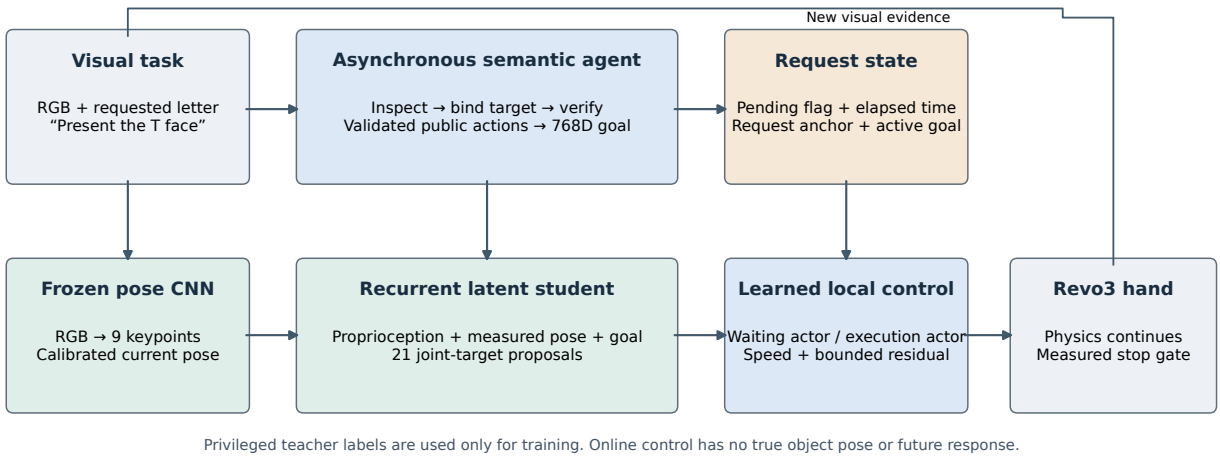


Figure S1. **Expanded system architecture.** The diagram shows privileged teacher supervision, the semantic-to-latent student interface, and closed-loop deployment. This schematic describes the cube implementation and does not report hardware outcomes.

C.1. Semantic and motor inputs

The motor policy receives 109 target-free proprioceptive/history features, seven measured pose values, and a 768-dimensional goal. Its goal projection has width 64 and a tanh nonlinearity. The inherited recurrent policy emits 21 normalized joint targets. The asset-specific goal quaternion is absent from the migrated student's inference inputs.

The hosted model route recorded in the manifests is `gpt-6-astra`. The local grounding backbone is Qwen3-VL-4B-Instruct. The REVEAL pathway uses public tool actions and a deterministic goal-bank lookup. The fine-tuned pathway predicts a continuous vector within the normalized skill span. Neither pathway receives the requested face's true pose. The local unreadable-view schedule is $(1, 3, 0, 2, 4, 5)$ with 90-step inspections. REVEAL chooses inspection actions of 10–90 steps. This scheduler difference is part of the complete-system comparison and a confound for attributing all effects to model weights.

The language-study archive contains exact task prompts and schedules. The route name alone does not identify immutable serving weights. **[TODO: Reconcile available serving metadata and sampling settings, and include full sanitized system prompts/tool schemas for each release in the anonymous package; state explicitly which provider revisions were not exposed.]**

C.2. Grounding supervision

The grounding model first uses 4,096 teacher-view training examples and 1,024 validation examples, then 4,800 student-view training examples and 2,400 validation examples with mixed replay. Each adaptation stage uses 240 updates. Training and validation collection seeds are disjoint from the formal evaluation seeds. Readability is a geometric/image-signal proxy. The objective combines readability prediction, readable-example skill classification, and cosine alignment to a canonical vector. We do not claim human-verified text recognition labels or an action-consistency objective in this release.

The fixed layout permits learned letter-to-color shortcuts. A held-out pose or trajectory alone does not remove that possibility. A convincing grounding test changes or hides the visual evidence while holding the instruction and physical state fixed, then checks whether the model's selected target changes appropriately.

C.3. Controller features and losses

The sensor branch uses the current proprioceptive vector, measured pose, student belief and proposal, request-anchor pose and action, fitting residual, binding flag, pending flag, and elapsed pending duration. Its normalized inputs are clipped to $[-8, 8]$. Goal features are centered and scaled using the fixed bank. The sensor, goal, and fusion widths are 192, 64, and 128. Each actor has 22 outputs: a speed logit and 21 residual values.

The waiting actor interpolates from current joint positions toward the proposal and adds a bounded residual. The execution actor applies a bounded residual to the proposal at full speed. Both are supervised from the teacher with residual regularization of 0.01. Waiting labels hold the request anchor. Execution labels follow the active goal. Paired state/goal examples and schedule-derived states are used during fitting. Their causal utility remains to be established by matched training ablations.

No raw RGB, true object pose, teacher action, future model response, or oracle semantic label is a deployed actor input. Vision affects the actor through estimated pose and recurrent belief. The reported stopping gate is measured-

pose based. Previously trained completion-evidence heads are inactive and cannot be credited for the current results.

D. Earlier System Comparisons

Semantic condition	Seed 1	Seed 2	Raw total
REVEAL	13/20	16/20	29/40
Fine-tuned VLM, bare	14/20	13/20	27/40
Rule VLM, bare	5/20	9/20	14/40

Table S2. Raw semantic successes by seed. Seed 1 is 183101 and seed 2 is 183102. The bare addendum contributes the second and third rows only.

Comparison on eligible trials	A-only	B-only	Exact p
REVEAL / bare fine-tuned	9	6	0.6072
REVEAL / bare rule	18	2	0.0004

Table S3. Exploratory two-sided exact McNemar tests. Both comparisons use 39 common eligible trials. Tests are exploratory and unadjusted.

Individual rate intervals use the Wilson method with nominal 95% coverage. Paired comparisons use trial identities and initial-state verification. They account for pairing but do not address all uncertainty from hosted model variability or future object distributions. A single two-seed cohort is insufficient to establish a stable ordering between systems separated by two or three raw successes.

E. Successor Student: A Separate Development Study

The existing cube results use a fixed student weights. Successor training is designed to ask whether extra DAgger with the actual frozen CNN sensor improves physical control beyond an equal budget with simulated pose noise. Both branches start from that same latent student, use training seed 183201, and run 400 updates with 16 environments, 45 rollout steps, and four fitting epochs. Learning rate is 3×10^{-4} with a new Adam state and zero weight decay. The teacher action fraction decays as $0.9 \cdot 0.95^{i}$. Each branch collects 288,000 transitions. The CNN and teacher remain frozen.

After training, the final model from each branch is compared with the released student and the frozen latent student. Every arm uses CNN observations and 40 fresh diagnostic trials from seeds 183221 and 183222. No intermediate model is selected from those trials. The CNN branch must exceed both latent references in reach count without increasing drops to pass the development admission rule. This gate is an engineering selection rule, not a significance claim.

A selected new student still requires separately collected controller data and compatibility validation. Only then may a new semantic release use reserved seeds 183301 and 183302. The original cube-v1 results remain unchanged. Both branches and all 160 diagnostic trials have completed with valid physics logs. The four reach counts are 39/40, 38/40, 39/40, and 40/40 for the released, frozen latent, extra simulated, and extra CNN policies. Drops are one, two, one, and zero. The CNN candidate passes the stated gate and has not been automatically deployed. Exact paired tests against the frozen latent and extra simulated policies give $p = 0.5$ and $p = 1.0$. Thus this is a promising development result, not an established perception-matching effect.

F. Additional Experiments

These studies retain the frozen task success criterion and do not replace the completed cube-v1 cohort. Service-eligible success is the primary metric. Cross-method comparisons remove the same service-affected slots from every arm. All raw records, including excluded physical drops, remain available. Forty scheduled slots do not necessarily yield forty eligible trials.

F.1. Completed independent replication

Seeds 183411 and 183412 each provide 20 trials per method. All 120 scheduled records pass the physical-log and initial-state checks. Slots 183411008, 183411015 and 183411017 have recorded REVEAL service errors and are removed from every arm for the matched comparison. The last includes a connection failure after approximately 126 seconds and a waiting-phase drop. Its exclusion follows the existing service rule. Raw success counts are 26/40, 29/40 and 22/40 for REVEAL, bare fine-tuned VLM and bare rule VLM. The common eligible counts appear below.

Method	Eligible success	Rate (%)
REVEAL	26/37	70.3
Fine-tuned VLM, bare	27/37	73.0
Rule VLM, bare	21/37	56.8

Table S4. Completed independent replication on 37 common service-eligible initial states from 40 scheduled trials per method. This supersedes the archived partial snapshot and remains separate from the original cube-v1 cohort.

The eleven eligible REVEAL failures comprise nine drops, one false stop and one abstention. The drops include one during startup, two during subsequent waiting and six outside waiting. Raw REVEAL counts additionally include one service-excluded drop. Against the fine-tuned VLM, REVEAL has five exclusive successes and six exclusive failures (exact paired $p = 1.0$). Against the rule VLM, the counts are eleven and six ($p = 0.332$). These exploratory comparisons

do not establish superiority. The earlier 18-slot snapshot remains archived for provenance but is no longer used for the follow-up table.

F.2. Completed visual interventions

The upper-level image changes while the low-level pose predictor continues to receive the original image. Lighting is fixed. The permuted condition assigns the next symbol in (T, E, R, L, S, H) to each color face and changes the requested symbol accordingly, preserving the scored physical target across conditions. The masked condition hides the lettering. Geometry-consistent image construction and the transmitted image hashes are audited. Seeds 183441 and 183442 provide 20 trials each, and physical initial-state hashes match across the three conditions.

Method	Identity	Permuted	Masked
REVEAL	32/40	32/40	0/39
Fine-tuned VLM, bare	30/40	16/40	3/40
Rule VLM, bare	23/40	23/40	0/40

Table S5. Completed upper-image interventions. Identity and permutation each use 40 trials. REVEAL masking has one service exclusion. For its paired masking contrast the identity numerator becomes 31/39. All conditions use reconstructed face images, not the untouched native-image comparison.

Fine-tuned success falls by 35 percentage points after remapping and to 7.5% after masking. Rule-based success is 57.5% under both identity and remapping, then zero after masking. Equal aggregate identity and remapped rates do not imply identical trial outcomes. REVEAL now has 32/40 successes under both identity and permutation, with zero of 39 eligible masked trials. This supports robustness to the tested permutation, without proving equivalence or identifying all fine-tuned successes as a fixed lookup.

Applied semantic bindings. We audit every applied, accepted call from the same 240 trials. The episode-file hashes match the archived results, and the last accepted skill agrees with the recorded terminal binding. Table S6 separates acquiring the correct target from completing it. Under permutation, the fine-tuned model first binds correctly in 25 trials and incorrectly in eleven, compared with 34 and four under identity. Nine of its eleven wrong-binding trials select the original color for the newly requested symbol. The rule model has six wrong-binding trials under each layout, with three following the old association after permutation. This pattern is consistent with reliance on training associations, but does not identify the model's internal mechanism.

No wrong-to-correct binding transition occurs in these bare-VLM cohorts. This is descriptive: trajectories, visibility, and opportunities for additional calls differ. Of fourteen fine-tuned false stops under permutation, ten end with a wrong binding and four with the correct binding but an independently incorrect final pose. Masking leads to no accepted binding in 36 fine-tuned trials and all forty rule-model trials. The three masked fine-tuned successes cannot establish letter grounding when that information is removed.

Model	Layout	Correct	Wrong	None	Old
Fine-tuned	Identity	34	4	2	–
Fine-tuned	Permuted	25	11	4	9
Fine-tuned	Masked	3	1	36	–
Rule	Identity	24	6	10	–
Rule	Permuted	25	6	9	3
Rule	Masked	0	0	40	–

Table S6. First accepted binding in each of 40 trials. Old counts trials with an accepted skill matching the original symbol–color association, defined only for permutation. These are the same trials as the success table, not new evaluations.

Identity versus permutation yields eighteen exclusive successes and four exclusive failures for the fine-tuned model (exact paired $p = 0.00434$), versus ten and ten for the rule model ($p = 1.0$). The final analysis applies Holm correction across six planned image contrasts, giving adjusted $p = 0.0130$ for the fine-tuned mapping contrast. The single deterministic remapping and reconstructed upper images limit generalization claims. REVEAL has five exclusive identity successes and five exclusive permutation successes ($p = 1.0$).

F.3. Q3: Initial planning and continued feedback

The two conditions share 40 initial states. Full feedback reuses Q2's 40 trials from seed 187102, so these trials are not an independent replication. The initial plan uses one observation to return a bounded action sequence. Full feedback uses updated observations and history on up to six calls. Both conditions use the same student, local controllers, stopping rule and time limit.

Protocol	Raw success	Eligible	Drops	Calls
One initial plan	6/40	6/39	0	40
Continued feedback	34/40	33/39	5	120

Table S7. Q3 feedback comparison. One initial-plan HTTP 503 timeout is excluded from both arms for paired analysis. Drops and calls include all scheduled trials.

Twenty-seven pairs succeed only with feedback, and none only with the initial plan (two-sided exact paired $p = 1.49 \times 10^{-8}$). Of 28 eligible initial-plan abstentions, 22 succeed with feedback. Among eligible feedback trials, 36 inspect, 32 first bind after inspection, and 31 successful trials use multiple calls. No trial corrects an executed wrong

binding. These counts describe how the agent gathers evidence. They do not tell us whether the target was initially visible or whether every first binding was correct.

Median logical duration is 8.17 seconds for initial planning and 22.0 seconds for feedback on the eligible panel. These durations include unsuccessful termination and must be interpreted alongside 28 initial-plan abstentions. Median measured call latency over raw requests is 8.12 and 5.33 seconds. The protocols differ in output structure and inference budget.

F.4. Three-seed controller evaluation

The controller ablation uses seeds 187102, 187103 and 187104, with 40 trials per seed and method. Seed 187102 provides the first 40 trials, and the other two seeds add 80. Q2 reports their combined counts as the paper's only formal online controller ablation. For learned slowing alone, we disable joint correction and keep the same initialization, data, sampling, optimizer and 2,400 updates. Only combined counts are available for the two additional seeds, so this table cannot support seed-level variance estimates or paired significance tests.

Waiting policy	187102	187103 + 187104
Continue	25 / 7 / 5 / 3	48 / 17 / 9 / 6
Fixed slowing	30 / 4 / 0 / 6	58 / 9 / 8 / 5
Learned slowing	28 / 1 / 7 / 4	54 / 6 / 7 / 13
Slowing + correction	34 / 4 / 1 / 1	70 / 1 / 3 / 6

Table S8. Outcome breakdown for the three-seed controller study. The two columns contain 40 and 80 trials per method, respectively. Each cell lists success / waiting drop / execution drop / other failure. Other failures are not further subdivided.

Q1's mapping study and Q3's image intervention use different trials. Their absolute rates need not match Q2, whose controller gains are measured within its own trial set. The initial-plan comparison shares the 40 full-feedback trials from seed 187102, as noted above.

F.5. Q3: Fresh visual evidence

The image intervention uses seeds 187201 and 187202, with 20 trials per seed and arm: 40 matched pairs, 80 trials in total. These are Q3 image-intervention seeds, not additional Q2 controller seeds. Condition order is reversed for the second seed. Both arms freeze the student, predictor, waiting and execution actors, physics, prompt, six-call cap and scoring rule. All 40 pairs are retained, with no service exclusions.

All 23 pairs with different success outcomes favor current images (two-sided exact paired $p = 2.3842 \times 10^{-7}$, unadjusted). Success counts for the two seeds are $(17, 15)$ with current images and $(5, 4)$ with repeated first frames. Wilson 95% intervals for the combined rates are 65.2–89.5% and 12.3–37.5%, respectively.

Upper RGB	Success	Drops	Calls	Tokens
Current image	32/40	3	112	209,547
Repeated first frame	9/40	4	98	161,082

Table S9. Q3 image intervention under a common six-call cap. All scheduled trials are included; realized inference cost is not matched.

Only the upper agent's RGB changes. The frozen arm sends exactly its first frame, including derived crops. Every request also saves the actual current frame. Pixel and file hashes are checked against transport records. All 40 initial states and first images match. Later live frames differ from the first on 72 fresh-arm and 58 frozen-arm requests. Both arms have 31 trials with multiple calls. Local perception, numeric pose feedback, history and motor control remain live in both arms.

Fresh-image outcomes comprise 32 successes, three drops, three false stops and two abstentions. Frozen-image outcomes comprise nine successes, four drops, three false stops and 24 abstentions. Median measured call latency is 6.25 versus 6.95 seconds, and median logical duration is 22.67 versus 25.62 seconds. Actual inference cost and trajectories differ, so this is not a matched-realized-compute test. Timing checks verify continuous logical control and delayed response application, without asserting wall-clock real-time simulation.

A post-hoc analysis checks whether each trial ends with a correct, incorrect or absent target binding. Current images give counts of 34, one and five. Repeated first frames give nine, three and 28. Correct binding alone is insufficient for success: two current-image trials bind correctly but fail the final pose check. These counts do not measure first-binding accuracy, initial target visibility or stability after stopping. The analysis script records the source hash and each trial's classification.

When correct bindings are acquired. A further post-hoc audit uses applied tool calls rather than terminal states alone. Both arms first declare the correct binding on call one in the same nine trials. With fresh images, an additional 25 trials first bind correctly on calls two through six: eight, nine, five, two and one trials, respectively. Of these 25, 23 subsequently succeed; the other two stop with an incorrect true pose. Repeated first frames produce no additional correct binding after call one. Six fresh-image trials and 31 repeated-frame trials never bind correctly. These are semantic declarations in recorded trajectories, not independently measured first visual-recognition events. All 210 request images were verified against transport hashes. Initial-target readability and finger-versus-back-face occlusion have not been independently annotated and are not reported as results.

F.6. Q3: Adaptive versus fixed inspection

This prospective addendum uses new seeds 187301 and 187302, with 20 trials per seed and arm (40 matched pairs; 80 formal trials). Two admission trials with seed 187291 are excluded from results. Both formal arms were run anew; the earlier fresh-image arm is not reused as a control. The order is adaptive then fixed for the first seed and fixed then adaptive for the second. All 80 outcomes are retained, with no service exclusions, replacement trials or extension after observing the result.

Intervention and controls. Both arms use the archived fresh-image cube release's semantic model (`gpt-6-astra`), system prompt, tool schema, pose estimator, latent student, waiting and execution actors, physical horizon and stopping rule. Each has the same six-call cap. In the adaptive arm, the model selects the color and duration of each `inspect_face` action, from 10 to 90 control steps. In the fixed arm, only these actions are replaced by a cyclic green, yellow, red, blue, magenta, cyan sequence, each lasting 90 steps at 30 Hz. The sweep index advances only on inspection. Non-inspection actions and target declarations are unchanged. Both arms retain the model's choice of whether to inspect, bind, verify, stop or abstain, so this isolates inspection direction/duration rather than removing all high-level autonomy.

The common feedback schema adds the last executed action in both arms, preventing an overridden fixed action from silently being treated as the model's proposed action. Numeric feedback and image history remain live; their realized values and subsequent trajectories differ. Initial physical-state and first-image hashes match for every pair. Runtime contracts, delivered images, applied actions and physical logs passed the final audit. Realized latency and compute are reported rather than assumed equal.

Inspection	Success	Bound	Unbound	Drops	Calls
Fixed sweep	31/40	32	7	1	126
Adaptive	38/40	39	1	0	117

Table S10. New Q3 inspection-policy cohort. Bound counts correct terminal bindings. Fixed inspection additionally has one wrong terminal binding; adaptive has none. Binding categories and terminal failure types overlap and must not be added together.

Successes by seed are $(18, 20)$ for adaptive inspection and $(16, 15)$ for the fixed sweep. Adaptive failures comprise one abstention and one false stop. Fixed failures comprise six abstentions, two false stops and one execution-stage drop. The drop phase describes timing, not its cause. Correct terminal bindings are 39/40 and 32/40, respectively.

Inference and acquisition cost. The prespecified primary endpoint is autonomous success, tested by exact two-sided paired McNemar. The paired table contains 29 shared successes, nine adaptive-only successes, two fixed-only successes and zero shared failures. The observed gain is 17.5 percentage points, but $p = 0.0654297$ does not reject equal success probabilities at $\alpha = 0.05$. Wilson 95% intervals for the individual rates are 83.5–98.6% and 62.5–87.7%; these are not paired-effect intervals. This is one primary test in a new study; binding and cost outcomes are descriptive secondary measures. The result is neither a statistically significant advantage nor evidence of equivalence.

First correct-binding call counts for calls one through six are $(8, 12, 12, 5, 2, 0)$ for adaptive and $(8, 5, 9, 6, 4, 0)$ for fixed inspection. One and eight trials, respectively, never bind correctly. The index includes the input producing the binding; strictly preceding observations equal the index minus one. These never-bound outcomes must accompany any conditional acquisition-time statistic; early failure cannot be interpreted as faster acquisition. The experiment records declared bindings, not an independent visibility annotation.

Adaptive inspection uses 117 calls and 247,383 tokens, versus 126 calls and 273,048 tokens for the fixed sweep. Median measured call latency is 4.57 versus 4.56 seconds; median logical episode duration is 17.83 versus 20.33 seconds, including failures. One prescribed sweep on one cube does not establish superiority to all reactive policies, and an equal maximum call count does not ensure equal realized compute. Different cohorts and hosted-service conditions also preclude attributing differences from older Q3 success rates to this intervention.

F.7. A recorded feedback sequence

Figure S2 shows a recorded pair from the initial-plan/continued-feedback comparison. Both protocols start from the same physical state and byte-identical initial image. The one-call plan abstains. The feedback agent inspects red, then binds T to red after the next observation. The measured gate stops at step 420, with independently scored orientation error 0.0812 radians and correct semantic binding. This is the lowest trial identity among eligible pairs with initial-plan abstention, feedback success, and an inspection action. It is a selected illustration, not a visibility annotation for the whole cohort or evidence of wrong-to-correct binding.

F.8. Latest native comparison and failure audit

Seeds 185451 and 185452 give 40 matched trials per method with no service exclusions. The REVEAL system gives 35 successes, three execution-phase drops, one wrong-binding false stop, and one abstention. The bare fine-tuned system gives 30 successes, nine execution drops, and one horizon failure. The bare rule system gives 17 successes, twelve drops, ten false stops, and one horizon failure. Its drops

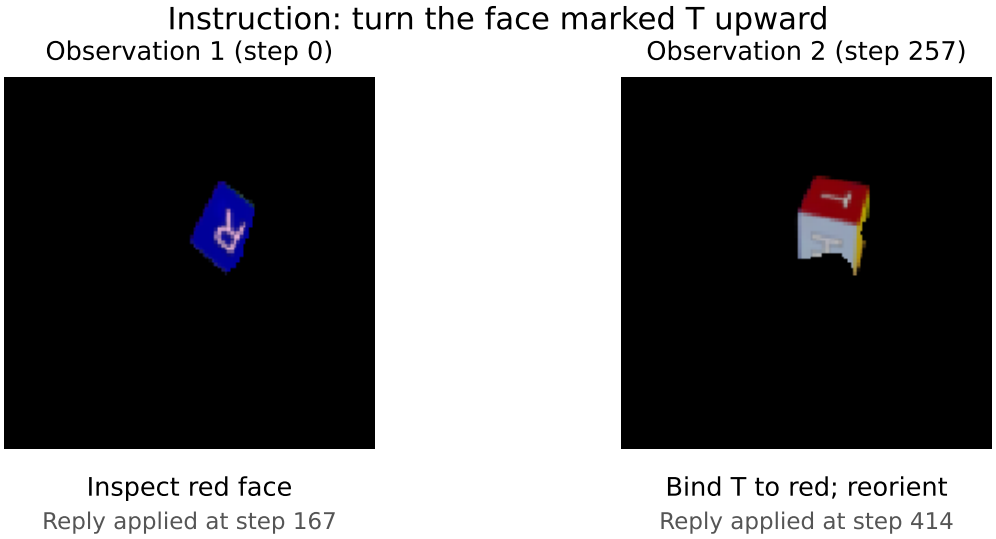


Figure S2. **Inspection followed by grounded execution.** Exact native model-input frames from a completed trial, enlarged without retouching. The requested letter is T. The first observation contains a visible R, and the second exposes T on the red face. Captions summarize recorded public actions. Control continues between request and reply at both observations. The two calls take 5.55 and 5.22 seconds. These are simulator images, not hardware observations.

include two during startup waiting. REVEAL makes 101 calls, fine-tuning 117, and the rule baseline 126. Median latencies are 4.349, 0.099 and 1.104 seconds.

The three REVEAL drops occur during inspection before final target binding. One is 59 control steps after the most recent applied inspection command, and the others are 36 and 24 steps after it. These intervals describe timing, not causal attribution. The false-stop episode binds R to cyan rather than blue. Its measured error to the declared target is 0.057 radians, while true requested-target error is 3.079 radians. It is a semantic binding failure and cannot simply be assigned to pose estimation. The abstention follows four inspections without a declared target. A favorable instantaneous pose during that trajectory does not meet the autonomous success rule.

F.9. Waiting-controller fitting

The waiting refit uses training collection seed 185301, validation seed 185302, and fitting seed 185311. These are training/validation seeds, not Q2 evaluation seeds. The inherited actor is updated 2,400 times. Actor batches allocate one half to pending states and one quarter each to unbound and bound execution states. Pending samples are uniform. Only the waiting outputs replace the deployed actor; execution remains unchanged. Formal Q2 evaluation uses the three seeds in Sec. F.4.

F.10. Why speed and joint correction are distinct

A scalar coefficient scales the commanded joint displacement toward the student proposal. It cannot independently change the command direction. Actual motion also depends on the joint servo and contact dynamics. To quantify that restriction, we use 17,755 pending states from 61 validation trajectory groups collected with seed 185302. These states are correlated and contain no formal evaluation trials. For normalized current position q, student proposal u, and the teacher holding action a^*, the best scalar action in hindsight is

$$\alpha^* = \mathrm{clip}\left(\frac{(a^* - q)^\top (u - q)}{\|u - q\|_2^2}, 0, 1\right). \tag{S1}$$

If $u = q$, the interpolation action is q for every coefficient. The resulting error is a lower bound for any scalar-speed predictor on that state. It uses privileged teacher labels for analysis and is not a deployable policy.

The full actor is below the scalar bound in 64.8% of these states and has a lower mean error than the bound. This supports the additional action space provided by joint correction. It does not show that every residual is useful or that lower action error necessarily improves closed-loop success. The three-seed Q2 evaluation tests that separate question.

F.11. Timing verification

The request logs verify that pending requests preserve continuous logical control and that replies are applied only after

Command design	Teacher-action MSE
Continue student	0.5953
Fixed slowdown	0.6095
Learned slowdown	0.4906
Learned slowdown + correction	0.3870
Best scalar in hindsight (bound)	0.4330

Table S11. Held-out waiting-action fit in normalized joint-command coordinates. Lower is better. The privileged bound is distinct from the four deployable conditions. These values are not task success rates.

their measured latency and actual completion. Replies arriving after episode termination are not applied. These checks verify the simulated delay protocol rather than wall-clock real-time operation. Q2 aggregate outcomes do not provide per-trial timing for the two additional seeds, so no timing result is extrapolated to all 120 trials.

G. Scope, Interpretation and Remaining Work

G.1. Separate model strength from control

The primary comparison uses REVEAL against bare VLM baselines. Its interpretation is system-level. To attribute gains to the upper-level model, compare all providers with identical control actors, timing conventions, stopping, and information access. The current Q2 table varies pending-request control within REVEAL; it does not supply this matched comparison across semantic models. Repeated live calls can produce different public plans even on matching initial states. Fixed-public-plan replay can isolate motor changes, but replay success must not be reported as a new live REVEAL success rate.

G.2. Test semantic and visual causality

The current Q2 study does not isolate the utility of explicit controller-goal input; no benefit is attributed to that pathway. Q1 reports semantic-layout sensitivity separately for the cube and die. These complete-system interventions do not extend the Q2 controller findings to other objects. No additional multi-object controller ablation is included in the current experiment scope.

G.3. Test agent feedback

Q3 now contains three separate cube cohorts: continued feedback, fresh images, and adaptive versus fixed inspection. The first changes call budgets; the second changes upper RGB under a common cap; the third changes inspection direction/duration and shows a favorable but statistically unresolved success difference. None isolates every history or verification component, equalizes realized compute, or establishes a mechanism across objects.

G.4. Separate perception accuracy and control benefit

The independent pose study below establishes frame-level accuracy differences between two supplied systems. To test downstream benefit, compare predictors with the same student and agent. Report false stops, drops and semantic success alongside independent pose error. Oracle pose estimates the residual motor gap. Only a second experiment should adapt the student to the new predictor. A lower pose error is useful evidence, but it does not guarantee a higher task success rate.

G.5. Objects and physical hand

Multi-object evaluation is confined to Q1 complete-system comparisons and Q4 physical-hand experiments. For Q1, the current scope is six objects with variation in geometry, inertia, contact affordances and semantic ambiguity. Q4 must identify the actual physical objects tested and report its coverage separately. Record object-specific students and training budgets. Balance the task set within each object and report each row before computing an object macro-average. Forty trials per arm per object is the default batch size. Hardware runs need calibrated independent scoring, documented resets, two sessions, and synchronized video and command logs. Each object requires its assets, calibration, action conventions and goal definitions.

H. Independent Pose Estimation under Hand Occlusion

This supporting perception study for Q4 evaluates fixed RGB and RGB-D estimators on a 58.5 mm lettered cube. Its RGB model differs from the frozen agent CNN, and its pose recalls do not update agent success counts.

Sampling and scoring. Three test seeds (31001–31003), four camera pitches ($0°$, $60°$, $90°$, $105°$), and 512 frames per group yield 6,144 observations in 48 environment clusters. RGB, RGB-D, their no-hand controls and perturbed-depth RGB-D produce 30,720 predictions. Collection records four geometry rejections; 6,139 admitted observations contain the full object.

Recall requires translation error < 10 mm and full rotation error $< 10°$, without cube-symmetry minimization. Invalid predictions count as failures. Secondary recall uses 5 mm and $5°$. ADD uses 6,144 nominal-cube surface points and threshold 0.1 diameter, not BOP average recall. Frozen prose used non-strict inequalities, but all evaluations consistently used strict thresholds. No samples lie on the primary boundaries.

Accuracy and uncertainty. RGB passes on 2,306/6,144 observations and RGB-D on 3,977/6,144, a difference of

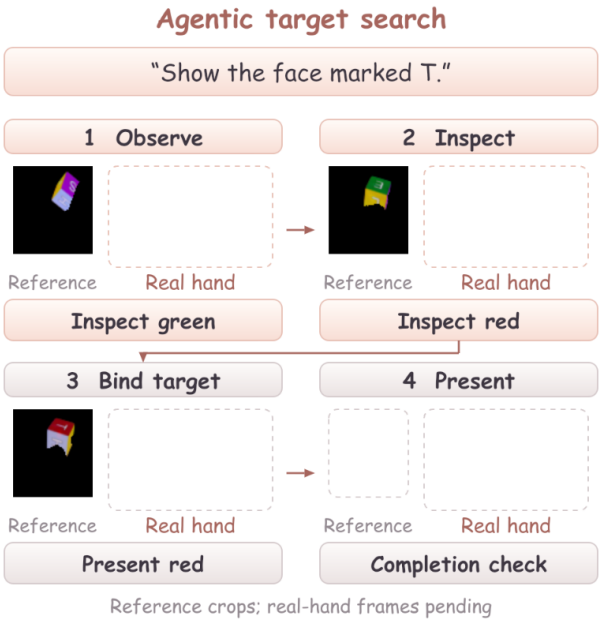


Figure S3. **Planned inspection and presentation sequence.** Reference crops illustrate observing, inspecting, binding, and presenting a requested face. Real-hand image slots remain empty. This is a protocol mockup, not a physical-robot result.

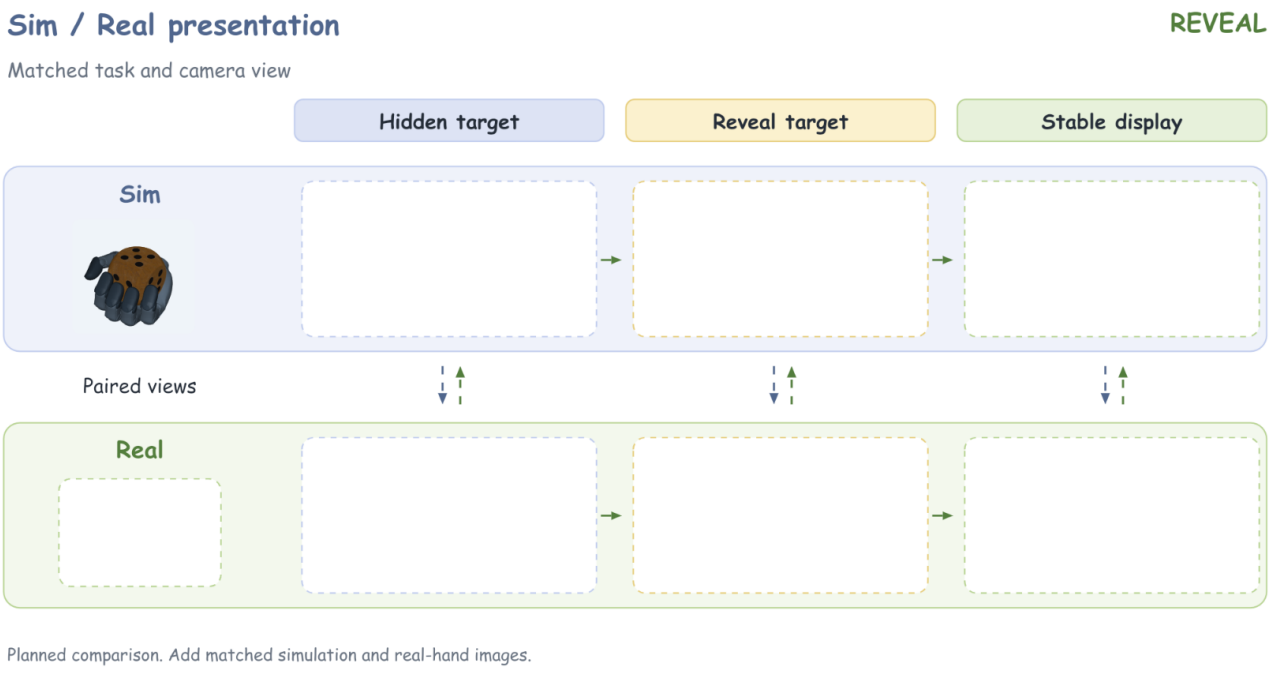


Figure S4. **Planned simulation/physical presentation comparison.** The layout pairs hidden-target, reveal, and stable-display stages. The reference simulation crop is illustrative; the remaining simulation and physical-hand slots await matched observations. No real-robot outcome is represented.

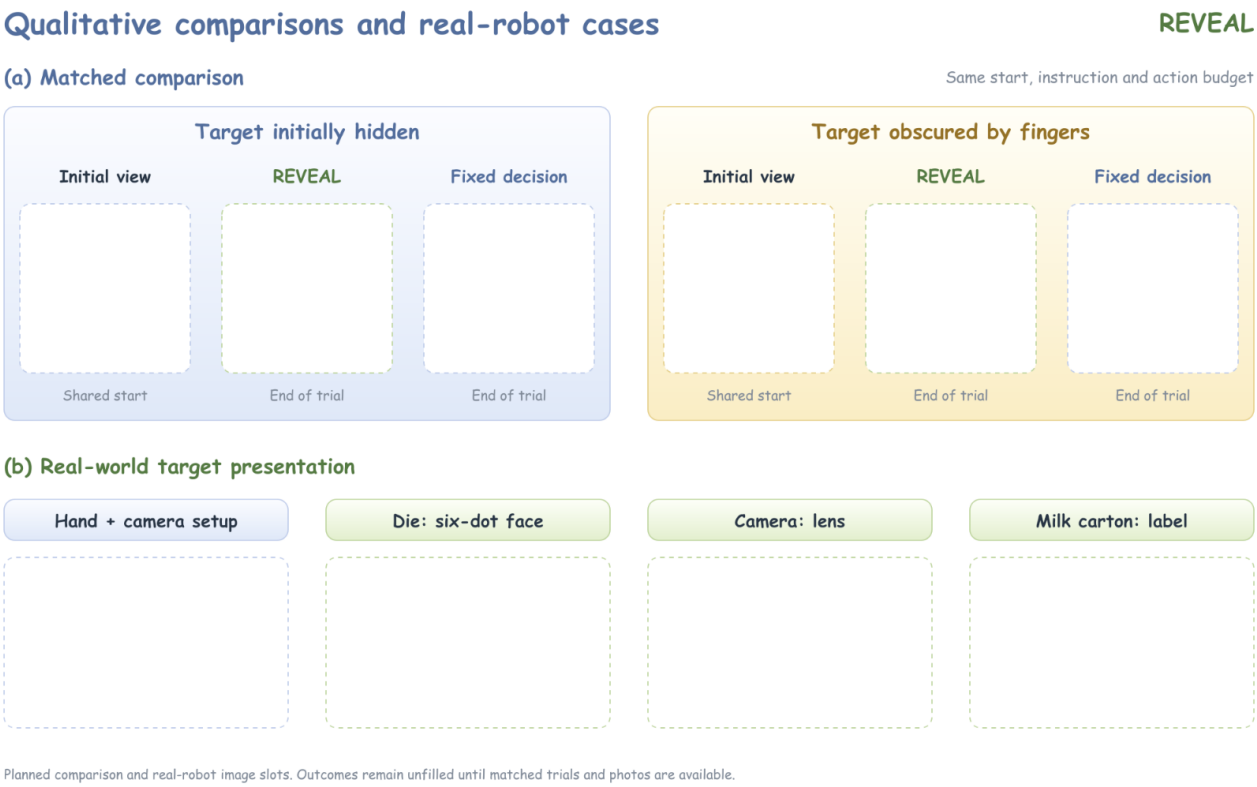


Figure S5. **Planned qualitative comparison and real-robot cases.** Initial/end-state and physical-robot image slots remain placeholders until corresponding observations are available. No completed trial is represented.

Hand occlusion	Frames	RGB	RGB-D
[0,25%)	2,891	66.17	97.61
[25,50%)	1,964	19.91	55.45
[50,75%)	1,198	0.17	5.51
[75,100%]	91	0.00	0.00
All	6,144	37.53	64.73

Table S12. Independent pose recall (%) under hand occlusion. All frames remain in the denominator, including invalid predictions. Occlusion is associated with viewpoint in this sampling design.

Metric	RGB	RGB-D
Recall, 10 mm / 10 deg. (%)	37.53	64.73
Recall, 5 mm / 5 deg. (%)	14.63	47.56
ADD recall (%)	38.35	71.48
Translation mean (mm)*	57.56	6.26
Rotation mean (deg.)*	12.28	10.82
Invalid predictions	69/6144	1/6144

Table S13. Secondary pose metrics. All recall denominators include invalid predictions. *Mean errors use valid predictions only and must be interpreted with the invalid counts. Threshold comparisons are strict.

27.20 percentage points. The reported 95% intervals are [36.60,38.48]% for RGB, [64.02,65.46]% for RGB-D, and [26.07,28.30] points for the paired difference. Intervals use 10,000 paired bootstrap draws of whole environment clusters, stratified by seed, retaining their views and temporal observations together. They describe sampling variation for this object and these test seeds, not variation across trained models or unseen objects. Table S12 shows that gains concentrate in light and moderate occlusion. At 50–75%, RGB-D recall is only 5.51%, and both systems fail all 91 observations at $\geq 75\%$.

Paired controls and sensitivity. Removing hand occlusion at matched poses gives recall of 76.90% for RGB and 99.89% for RGB-D. The overall occlusion-associated loss is therefore 39.37 versus 35.16 points, with a paired reduction of 4.22 points [2.77,5.62]. This reduction is not uniform: at 50–75% occlusion, the loss is 72.20 versus 94.49 points. The higher unoccluded RGB-D ceiling does not imply robustness when most of the object is hidden. The no-hand RGB render retains soft Gaussian-splat boundaries, whereas the occluded image uses a hard visible-object mask, which also limits

this control. Restricting to fully framed observations gives 37.56% versus 64.78%, similar to the full cohort.

Adding 1 mm Gaussian ray-depth noise and dropping 20% of points yields RGB-D recall of 3,973/6,144 (64.66%). This is a sensitivity check, not a calibrated physical depth-camera error model. Per-seed RGB/RGB-D recall is 37.30/64.89%, 37.70/64.36%, and 37.60/64.94%. Viewpoint and occlusion covary: recall at nominal view is 76.04/99.80%, versus 0.85/4.23% at 105°.

Training and interpretation. The RGB system uses development-selected letterboxing and its existing PnP decoder. RGB-D uses the supplied epoch-1812 weights. Training histories and camera adaptation differ, so this is not a matched-budget architecture or depth-input ablation. RGB-D training manifests record 20,000 samples from seeds 43–50, validation uses seed 2027, and the test seeds do not overlap. Only nine training entries have more than 50% occlusion. This post-hoc manifest observation is consistent with the severe-occlusion failure but is not causal evidence, and the full training imagery has not been re-audited. The old RGB training-data lineage is incomplete.

The evaluation uses simulator pose labels, ideal target segmentation, and virtual co-located depth. Real-camera calibration and Gaussian-splat appearance assets are part of the setup, so it should not be described as using no real-world information. The evidence supports improved simulated pose accuracy over this RGB deployment, with a clear severe-occlusion limit. It does not yet establish improved agent completion or real-camera accuracy. The earlier 39/40 original pose-goal versus 35/40 latent-student physical-reach diagnostic used the frozen agent CNN and is a separate experiment.

Evidence provenance. The supplied integrity report records 129 unchanged frozen files, shared finite ground truth, and pose-error recomputation within 2.28×10^{-13} mm and 1.01×10^{-11} degrees. Paper integration verifies file hashes and re-sums twelve shard counts. Bootstrap intervals and per-frame checks retain delivery provenance. Protocol/model hashes are archived; the final reproduction package still needs sanitized per-frame records.

I. Q1: Native-Layout Dice Extension

The die is a wooden object with one through six native textured pips, rendered at 512×512 with the hand. Three methods each run 40 matched trials over seeds 186301 and 186302. The independent audit verifies all 120 records, canonical targets, paired reset states, response latency and baseline isolation. No trial is service-excluded.

System	Success	%	Drops	Calls
REVEAL	24/40	60.0	15	63
Fine-tuned VLM, bare	20/40	50.0	8	141
Rule VLM, bare	17/40	42.5	7	154

Table S14. Completed native-layout die comparison. No service exclusions. The independent success definition is shared; internal control and stop modules differ by system.

Object-specific motor interface. The released feedforward actor uses ten frames of normalized input. Its goal-related channels are replaced by a learned reconstruction of eight conditioning features. For each frame, a 768-dimensional latent is normalized and encoded to 32 dimensions. The adapter combines it with the four measured quaternion components and their outer product, then applies layers of widths 256, 128, and 8. The original actor and observation normalization remain frozen. Inference receives physical history of shape 10×116 and continuous latent history of shape 10×768. It receives no true target quaternion, target index or table-based reconstruction of its conditioning features.

Perception, control and stop policy. The die uses the released simulator-pose noise, bias, failure, jitter and delay model. There is no image pose estimator in these trials. All upper methods receive the same RGB and image-only crop. The physical horizon is 120 seconds at 20 Hz and the stop gate is checked at 3 Hz. A reply cannot be applied before $\lceil 20L \rceil$ steps, while control continues. The drop boundary and independent full-quaternion success tolerance remain 0.24 metres and 0.16 radians.

REVEAL uses one object-specific maintenance actor during pending and ordinary execution. Its calibrated completion model learns a bounded correction to measured angular error and an uncertainty term. It accepts when the calibrated upper bound is at most 0.16 radians. Once a correct count is declared and the state is not inspection, this gate may stop even with a further request pending. The bare baselines use no learned maintenance or completion module and stop on the measured-angle criterion only when no request is pending. Thus these are complete-system comparisons, not isolated tests of the semantic models or a replication of the cube's shared measured gate.

REVEAL's 95% Wilson interval is [44.6, 73.7]%, versus [35.2, 64.8]% for fine-tuning and [28.5, 57.8]% for rules. REVEAL has ten exclusive successes against six for fine-tuning (exact paired $p = 0.4545$), and eleven against four for rules ($p = 0.1185$). These exploratory comparisons do not establish superiority. REVEAL has twelve startup drops, three execution drops and one false stop. Fine-tuning has eight execution drops, seven false stops and five horizon

System	Identity	Permuted
REVEAL	28/40	26/40
Fine-tuned VLM, bare	23/40	0/40
Rule VLM, bare	13/40	1/40

Table S15. Dice visual-layout transfer. All 240 scheduled trials are service-eligible. Forty reset states and physical goals are matched across all six cells.

failures. Rules have seven startup drops, six false stops and ten horizon failures. Median request latencies are 5.63, 0.124 and 1.16 seconds, respectively.

J. Q1: Completed Dice Visual-Layout Study

The six cells use identity or a 120-degree visual-mesh rotation about a body diagonal, crossed with three methods. Each cell contains 40 scheduled trials over seeds 186701 and 186702, with matched physical targets and initial states. Collision, inertia and formal solver settings are unchanged. The mesh has slight asymmetries, so this is a visual-layout intervention rather than a pure texture edit. Independent deterministic diagnostics produce identical 32-step state/action traces for native, identity and permuted layers. The formal solver retains its original settings.

Input and motor contract. Printed counts are separated from fixed physical motor IDs. Skills 1–6 display counts $[1, 2, 3, 4, 5, 6]$ in identity and $[2, 4, 6, 1, 3, 5]$ after rotation. This table is used only to construct task instructions and score outcomes. Models receive no layout table, and their commands are not inversely remapped. The learned VLM's continuous 768-dimensional output is applied unchanged. Both layouts share an explicit skill-ID prompt that differs from the native-main prompt, making identity a control for that common change. The VLM retains its native-layout-trained weights. This measures zero-shot layout transfer, not performance after training on permuted examples.

REVEAL receives updated RGB, action/observation history and noisy feedback and can probe skills. The bare VLMs receive only current RGB and instruction, with their original fallback inspection schedule. From a single view, a new printed-count-to-motor-ID association may be ambiguous without a stable face reference. The comparison therefore measures complete systems, not a controlled effect of model scale or fine-tuning alone. Dice maintenance and stopping modules remain as documented above. The shared primary success predicate is unchanged.

Results and paired uncertainty. REVEAL completes 28/40 identity trials and 26/40 permuted trials. Their 95% Wilson intervals are [54.6, 81.9]% and [49.5, 77.9]%. Fine-tuning changes from 23/40 to 0/40, and rules from 13/40 to 1/40. Four paired states succeed only for identity REVEAL and two only for permuted REVEAL. The absence of a detected difference does not establish equivalence or robustness to other layouts. Table S16 reports all seven exploratory paired contrasts without multiple-testing adjustment. The native-main results are a separate cohort and must not be pooled with this study.

Contrast A versus B	A only	B only	Exact p
I: REVEAL vs. fine-tuned	11	6	0.3323
I: REVEAL vs. rule	18	3	0.0015
P: REVEAL vs. fine-tuned	26	0	2.98×10^{-8}
P: REVEAL vs. rule	25	0	5.96×10^{-8}
REVEAL: I vs. P	4	2	0.6875
Fine-tuned: I vs. P	23	0	2.38×10^{-7}
Rule: I vs. P	13	1	0.0018

Table S16. Exploratory exact two-sided McNemar tests on 40 paired states. Tests are unadjusted. I and P denote identity and permuted layouts.

Binding and control failures. Permuted REVEAL has 27 trials with a correct binding to a skill ID different from the printed count, and 26 complete successfully. It applies inspection in 32 trials and first binds after inspection in 27. No wrong-to-correct binding transition is observed. Its failures comprise seven startup drops, one execution drop, five abstentions and one wrong-binding stop. These traces describe evidence gathering before commitment, not recovery from an already executed incorrect binding.

The permuted fine-tuned VLM has 31 wrong-binding stops, five horizons and four execution drops. All 31 wrong-binding stops report the native same-number skill. Code review confirms that RGB and instruction enter the backbone and learned latent head, with no service-side assignment from requested count to skill. In identity, it binds correctly in 36 trials but succeeds in 23, illustrating a separate control/stop limitation. The permuted rule model has 32 wrong-binding stops, four horizons, three drops and one success; 24 of its wrong-binding stops select the native same-number skill. These post-hoc diagnostics are consistent with reliance on native associations. They do not establish memorization as an exclusive cause, especially given the baseline's limited observation history.

Audit and scope. The independent audit checks all 240 trials, 69 frozen source files, matched resets, model-request image hashes, delayed response application, continuous VLM latents and canonical-quaternion scoring. A separate SciPy calculation verifies all six Wilson intervals and seven exact paired tests. No service exclusions, outcome replacements or post-test model changes occur. Trace diagnostics were



Figure S6. A recorded permuted-layout episode requesting two dots upward. Left: probing physical skill 3 exposes six dots. Right: probing skill 1 exposes two dots. REVEAL then binds the target to skill 1 and stops with an independently measured orientation error of 0.025 radians. Images are the actual model inputs. This first scheduled trial was selected after its outcome and illustrates one successful sequence, not a reliability estimate.

developed while cells were completing and are explicitly post-hoc. Earlier invalid digit-overlay attempts are excluded and preserved separately. One changed layout, one native-trained local model and object-specific motor weights do not establish broad semantic or geometric generalization.

K. Four Additional Q1 Objects

Duck, hedgehog, pig and milk carton each contribute six cells of forty raw trials, totaling 960. Each object uses its own frozen student and adaptation, native marked RGB for semantic decisions, and noisy simulator pose for motor control. Bare VLM baselines omit the added control modules. We preserve every recorded outcome. Two historical duck service failures and one hedgehog service failure affect only original-mapping cells. Common service-eligible panels remove the corresponding slots from all three methods. All changed-mapping cells retain forty trials.

The full audit covers all 960 episodes, 320 within-mapping initial-image groups and 160 physical-state groups. Seventeen image comparisons differ in one color channel of one pixel by one 8-bit level. Saved model-request pixels and delivered-input hashes match strictly. One duck original-mapping REVEAL reset (slot 7) differs from the other conditions in one quaternion component by 2.98×10^{-8}. The normalized angular difference is 5.77×10^{-8} radians, with identical initial image pixels. We retain this explicitly reviewed numerical exception rather than claiming bitwise reset equality. All other recorded state components match exactly.

As a sensitivity analysis, removing slot 7 from every method gives original-mapping successes of 24/37 for REVEAL, 13/37 for fine-tuning and 15/37 for rules. Removing it from the changed mapping gives 23/39, 2/39 and 9/39. The ordering is unchanged. The main table retains the frozen service-exclusion rule. This post-run sensitivity check does not replace outcomes or change the success definition.

L. Q1: Language and Wording Coverage Across Three Objects

Three separate cohorts contribute 2,160 recorded trials: cube, duck and hedgehog, each with three systems, two languages, two wording templates per language and 60 trials per cell. Cube seeds are 210501–210503, duck seeds 210601–210603, and hedgehog seeds 210701–210703, twenty trials per seed. Each of six goals occurs ten times per cell. These are evaluation blocks, not independent training seeds. Within each object, the same 60 reset seeds and physical goals are reused across methods and instruction variants. Four prompt variants must not be pooled as 240 independent physical tasks.

Intervention and deployment. Mapping, task criterion and object-specific weights remain fixed within each system. Cube retains its visual pose estimator, goal bank, latent student, 30-Hz control and measured-pose stopping; REVEAL uses its paper-selected waiting and execution actors. Duck and hedgehog use their own latent students, 20-Hz control and noisy simulator-pose feedback; REVEAL additionally uses object-specific learned maintenance and completion correction. Bare local baselines omit these added actors. All systems use the independently scored 0.16-radian tolerance. Upper RGB is 120-by-120 for cube and 512-by-512 for duck

System	English Standard	English Rewrite	Chinese Standard	Chinese Rewrite
Cube: $n = 60, 60, 60, 60$				
REVEAL	42/60	47/60	48/60	36/60
Fine-tuned VLM	42/60	44/60	46/60	47/60
Rule VLM	28/60	27/60	18/60	19/60
Duck: $n = 60, 60, 60, 60$				
REVEAL	41/60	40/60	41/60	46/60
Fine-tuned VLM	23/60	18/60	19/60	21/60
Rule VLM	30/60	15/60	20/60	19/60
Hedgehog: $n = 60, 60, 60, 60$				
REVEAL	44/60	46/60	39/60	46/60
Fine-tuned VLM	23/60	19/60	21/60	15/60
Rule VLM	28/60	23/60	21/60	22/60

Table S17. Raw completion counts over all scheduled trials, retaining infrastructure failures in the denominator. These are separate object-specific cohorts; the four instruction variants reuse physical tasks.

System	English Standard	English Rewrite	Chinese Standard	Chinese Rewrite
Cube				
REVEAL	4	1	3	7
Fine-tuned VLM	0	0	0	0
Rule VLM	0	0	0	0
Duck				
REVEAL	0	1	0	0
Fine-tuned VLM	0	0	0	0
Rule VLM	0	0	0	0
Hedgehog				
REVEAL	0	0	8	0
Fine-tuned VLM	0	0	0	0
Rule VLM	0	0	0	0

Table S18. Service-excluded trial counts, each out of 60 scheduled slots. All raw records are retained.

System	English Standard	English Rewrite	Chinese Standard	Chinese Rewrite
Cube: $n = 56, 59, 57, 53$				
REVEAL	42/56	47/59	48/57	36/53
Fine-tuned VLM	40/56	43/59	43/57	43/53
Rule VLM	27/56	27/59	18/57	17/53
Duck: $n = 56, 57, 57, 56$				
REVEAL	39/56	40/57	40/57	43/56
Fine-tuned VLM	21/56	17/57	18/57	19/56
Rule VLM	29/56	15/57	20/57	19/56
Hedgehog: $n = 52, 52, 46, 55$				
REVEAL	38/52	40/52	36/46	43/55
Fine-tuned VLM	19/52	15/52	17/46	13/55
Rule VLM	22/52	18/52	16/46	19/55

Table S19. Strict pairing sensitivity: common service-eligible slots additionally require identical initial-state and first-image hashes across all three methods. Numerators are successes. This is separate from the main service-eligible estimand.

and hedgehog. The English-trained fine-tuned models receive the requested instructions without Chinese retraining. The rule baseline additionally uses localized task/output instructions for Chinese conditions. This is deployed-system language coverage, not an isolated change to the user sentence for every backend, a shared-policy transfer test, or a causal test of fine-tuning. Exact prompts are retained in the evidence bundle.

Service eligibility and raw records. The main-paper language table reports completion conditional on service availability. Within each object and instruction condition, the union of infrastructure-affected slots is removed from all three methods, including any successful baseline outcome at that slot. Thus each row comparison uses identical physical tasks and the same denominator. These rates do not measure end-to-end service reliability. This reporting choice was made after inspecting the raw outcomes; it is not described as preregistered. Table S17 retains every scheduled slot and Table S18 lists exclusions. No excluded slot is replaced; continuations visit only unvisited slots. Physical drops, incorrect stopping, abstention and task-horizon failures remain in the denominator. Infrastructure faults include TLS/proxy failures and server overload, not model grounding or control errors.

Continuation and availability. Service failures interrupted the hedgehog queue. Only unvisited tasks resumed, using separate workers and verified direct HTTPS to the same endpoint after a further SSL failure. Model, weights, prompts, seeds, scoring and one-attempt transport policy remained fixed. All prior outcomes, errors, actual latencies and scheduler changes are retained in the archive.

State pairing and image sensitivity. Cube and duck initial-state hashes match exactly across methods and instruction conditions. Four hedgehog slots differ only in object-quaternion components by at most 5.96×10^{-8}; all other recorded state components match exactly. A post-run numerical reset audit uses a 10^{-6} component tolerance, and the strict sensitivity below removes any nonidentical hashes. Cube first-image hashes match exactly. Some duck and hedgehog renders differ at the pixel level; these do not alter the service-eligible main denominator. Table S19 additionally requires identical initial-state and first-image hashes across all three methods in a condition. All strict paired tests are retained in the data export. The duck differences involve six slots, at most two RGB pixels and one channel level per pixel in 512-by-512 images. Hedgehog image differences involve 12 slots, at most two RGB pixels and one channel level in the same resolution.

System	EN vs. ZH Standard	EN vs. ZH Rewrite	Standard vs. rewrite English	Standard vs. rewrite Chinese
Cube				
REVEAL	1.0000	1.0000	1.0000	0.7637
Fine-tuned VLM	1.0000	1.0000	1.0000	1.0000
Rule VLM	0.4529	1.0000	1.0000	1.0000
Duck				
REVEAL	1.0000	1.0000	1.0000	1.0000
Fine-tuned VLM	1.0000	1.0000	1.0000	1.0000
Rule VLM	1.0000	1.0000	0.0936	1.0000
Hedgehog				
REVEAL	1.0000	1.0000	1.0000	1.0000
Fine-tuned VLM	1.0000	1.0000	1.0000	1.0000
Rule VLM	1.0000	1.0000	1.0000	1.0000

Table S20. Exact paired language/wording tests, Holm-adjusted jointly over 36 contrasts across three objects. Non-significance does not establish equivalence.

Paired inference. Two-sided exact McNemar tests use discordant pairs. Method tests use the same common three-method service-eligible panels as the main table. A language/wording contrast excludes the union of service-affected slots in its two arms; its denominator can differ from a main-table denominator. Holm correction is applied jointly across three objects: one family of 36 language/wording contrasts and a separate family of 36 method contrasts. The strict state-and-image-matched sensitivity applies the same corrections separately. Tables S20 and S21 report all main adjusted tests. No language/wording contrast survives Holm correction. 20 of 36 method contrasts survive correction. Strict state-and-image-matched sensitivity changes the 0.05 significance decision in 0 of the 72 contrasts; all discordant counts and adjusted values are archived. These tests do not isolate model-only or controller-only effects.

Scope and evidence. Two languages and two templates per language, with six fixed targets and one mapping per object, do not establish broad multilingual generalization or equivalence. Earlier misconfigured runs are excluded; historical native/mapping cohorts remain separate. The sanitized archive retains all 2,160 outcomes, 72 main and 72 sensitivity comparisons, prompts, schedules and checksums.

M. Reproducibility

The archived evidence retains outcomes, goals, stopping errors, service exclusions and latency for the audited cohorts. The three-seed controller aggregate is maintained separately from the historical trial records.

[TODO: Complete anonymous code and asset licenses.]

Condition	REVEAL–FT	REVEAL–Rule	FT–Rule
Cube			
EN standard	1.0000	0.1220	0.1487
EN rewrite	1.0000	0.0102	0.0424
ZH standard	1.0000	< 0.0001	< 0.0001
ZH rewrite	1.0000	0.0017	< 0.0001
Duck			
EN standard	0.0258	0.3458	1.0000
EN rewrite	0.0043	0.0011	1.0000
ZH standard	0.0043	0.0102	1.0000
ZH rewrite	< 0.0001	0.0001	1.0000
Hedgehog			
EN standard	0.0025	0.1120	1.0000
EN rewrite	< 0.0001	0.0004	1.0000
ZH standard	0.0013	0.0041	1.0000
ZH rewrite	< 0.0001	0.0011	1.0000

Table S21. Exact paired method tests on the main common service-eligible panels, Holm-adjusted over a separate family of 36 contrasts. FT denotes fine-tuning. Strict pairing sensitivity tests are archived.